

Project Assignment

Building a Robust MLOps Pipeline

Course: ID5003W – Industrial AI At Scale

Table of Contents

Abstract	3
1 Introduction	4
1.1 Problem	4
1.2 Motivation	4
1.3 Objectives	5
1.4 Scope and Assumptions	7
1.5 Summary	7
2 System Architecture and Design	8
2.1 Architectural Overview	8
2.2 Deployment Topology	12
2.3 Component Responsibilities	12
2.4 Data and Artifact Management	14
2.5 Control and Data Flows (Training/Data Pipeline)	15
2.6 Networking and Ports	22
2.7 Configuration and Secrets	23
2.8 Reliability and Fault Tolerance	23
2.9 Observability and Diagnostics	24
2.10 Security and Privacy	24
2.11 Performance and Scalability	24
2.12 Reproducibility and Lineage	25
2.13 Design Trade -offs	25
2.14 Risks and Mitigations	25
2.15 Extensibility	25
2.16 Summary	26
3 Tools and Resources	27
3.1 Compute & Operating Environment	27
3.2 Containerization & Orchestration	27
3.3 Source Control & Data Versioning	28
3.4 Distributed Data Processing & Feature Engineering	28
3.5 Model Development & Training	29
3.6 Experiment Tracking & Model Registry	29
3.7 Artifact Storage	30

3.8 Metadata Store	30
3.9 Model Serving	30
3.10 Monitoring & Drift Detection	31
3.11 Notebooks & Exploratory Analysis	32
3.12 Environment & Dependency Management	32
3.13 Logging, Observability, and Health	32
3.14 Testing & Validation Utilities	33
3.15 Dataset	33
3.16 Configuration & Secrets	33
3.17 Ports, Storage, and Performance	34
3.18 Limitations	35
4 Implementation Details	36
4.1 Environment & Project Scaffolding	36
4.2 Data Versioning with DVC (Inputs & Provenance)	36
4.3 Preprocessing & Feature Engineering (Spark)	37
4.4 Training & Hyperparameter Search (Spark MLlib + TVS)	38
4.5 Experiment Tracking & Registry (MLflow + Postgres + MinIO)	39
4.6 Serving API (FastAPI/Uvicorn) & Inference Logging	39
4.7 Drift Detection & Automated Retraining (PSI)	40
4.8 Container Images & Dependency Management	40
4.9 Configuration & Runtime Topology (Compose)	41
4.10 Integration Tests, Smoke Tests, and Evidence	42
4.11 Operational Behavior & Error Handling	43
4.12 Security, Privacy, and Ethics	43
4.13 Limitations and Future Extensions	43
5 Result and Analysis	45
5.1 Data and experimental context.....	45
5.2 Model configurations	45
5.3 Headline performance	46
5.4 Error analysis (confusion matrices).....	47
5.5 Feature Importance	49
5.6 Resource behavior and efficiency.....	50
5.7 API validation and tests.....	50
5.8 Future-Proofing and Robustness.....	51

5.9 Resource Utilization 52

5.10 Discussion 58

5.11 Recommended Scheme (with Justification)..... 61

5.12 Reproducibility and artifacts 63

5.13 Limitations and threats to validity..... 63

5.14 Summary 63

6 Conclusion & Future work..... 64

6.1 What we built—end to end..... 64

6.2 Technical Outcome & Empirical insights..... 65

6.3 Reliability, governance, and safety..... 65

6.4 Limitations 66

6.5 Future work 66

6.7 Summar 67

7 Reference68

Abstract

This work develops a fully reproducible, end-to-end MLOps system for Titanic survival prediction, engineered to run locally on Windows 10/11 with WSL2 and orchestrated via Docker Compose. The pipeline enforces data governance by tracking raw data with Data Version Control (DVC) and excluding raw files from Git history, while Apache Spark performs scalable preprocessing and feature engineering to produce a consistent Parquet corpus. Model development is conducted in Spark MLlib with automated hyper-parameter optimization (Train-Validation Split grids over Random Forest, Gradient-Boosted Trees, and Logistic Regression). Each experiment logs parameters, metrics, and diagnostic artifacts (confusion matrix; feature importances or coefficients) to MLflow with a PostgreSQL backend for metadata and MinIO (S3) for artifact storage. The best run is auto-registered in the MLflow Model Registry and automatically promoted to *Production*, providing a verifiable lineage from data to deployable model. Deployment is realized as a FastAPI microservice that resolves the current *Production* model from the registry, exposes health and prediction endpoints, and appends anonymized feature vectors to an inference log. Population Stability Index (PSI)-based drift monitoring compares live feature distributions against a stored baseline; when drift exceeds a configurable threshold and minimum traffic criteria are met, an automated retrainer re-executes the training pipeline, re-evaluates candidates, and promotes a superior model, thereby closing the loop from monitoring to lifecycle management. Collectively, the design demonstrates a rigorous, auditable, and automation-ready MLOps workflow that satisfies the project requirements while adding production-grade enhancements in tracking, registry governance, and drift-aware continuous improvement.

Keywords: MLOps, MLflow, DVC, SparkML

Code Repo Link: <https://github.com/bineshjose/ch24m521>

Chapter 1 : Introduction

The last decade has seen machine learning (ML) move from isolated notebooks to always-on services that make decisions in real time. As this transition has accelerated, the discipline of MLOps—the application of software engineering and DevOps principles to ML—has become essential. Effective MLOps demands much more than a high AUC on a static test set: it requires reproducible data pipelines, governed experiment tracking, model registries and stage gates, reliable serving, monitoring for drift, and safe retraining workflows. This report presents an end-to-end, production-style MLOps system built around the canonical Titanic survival prediction task. While the dataset is intentionally compact, the system is engineered to reflect the realities of operating ML in production: versioned data, distributed preprocessing and training, registry-based deployments, health checks, drift detection, and automated retraining—all packaged in containers and orchestrated to run repeatably on a laptop and scale to a cluster.

1.1 Problem

Predicting passenger survival on the RMS Titanic from tabular demographic and ticketing attributes (e.g., *Pclass*, *Sex*, *Age*, *Fare*, *Embarked*, *SibSp*, *Parch*) is a canonical classification task. While the dataset is modest in size, the engineering challenge we address is not the pure predictive problem but the full lifecycle: how to build a reproducible, auditable, and continuously improving machine-learning system—from data ingestion and preprocessing to training, evaluation, model governance, deployment, online monitoring, and automated retraining.

Formally, given a feature vector $x \in \mathbb{R}^d$ describing a passenger, the objective is to estimate $p(y=1|x)$, where $y \in \{0,1\}$ indicates survival. However, the real-world constraints—traceable data versions, consistent transformations, experiment tracking, registry management, drift detection, and safe promotion of models to production—are what elevate this from a classroom model to a production-grade MLOps system.

1.2 Motivation

1. **Reproducibility and governance.** Many ML efforts fail to reproduce results months later due to untracked data or mutable preprocessing. We explicitly

version raw data with DVC and codify all transformations in Apache Spark, ensuring that any result can be regenerated.

2. **Scalability of practice, not just data.** Even if Titanic is small, the patterns we introduce—Spark pipelines, containerized services, MLflow tracking/registry, MinIO (S3) artifact storage, and FastAPI serving—are scalable templates. The same blueprint applies to larger datasets and teams.
3. **Operational excellence.** Models degrade in the wild due to data drift and concept shift. We integrate Population Stability Index (PSI) monitoring and an automated retrainer loop to preserve performance without manual babysitting.
4. **Local-first, platform-agnostic deployment.** The solution runs entirely on Windows 10/11 with WSL2 and Docker Desktop, making it easy to evaluate, grade, and extend—while mirroring the architecture of cloud-native systems (registry, artifact store, tracking DB).
5. **Assessment alignment + real-world readiness.** The project satisfies each rubric requirement (data pipeline & versioning, distributed training with HPO, experiment tracking with artifacts, model registry with stage transitions, drift & auto-retrain) and adds professional touches (reload endpoint, diagrams, health checks, inference logging).

1.3 Objectives

We articulate objectives across technical, operational, and evaluation dimensions. Each is concrete and verifiable.

1) Data & Preprocessing Objectives

- **O1.1 – Data versioning:** Track raw data via DVC pointers (not raw files in Git); configure MinIO as an optional remote.

Evidence: Presence of data/titanic.csv.dvc, .dvc/config, and successful dvc push/pull.

- **O1.2 – Deterministic preprocessing:** Implement a Spark pipeline that imputes numerics, encodes categoricals, assembles features, and emits Parquet (data/processed/titanic.parquet) for training and evaluation.

Success criterion: dvc repro re-generates identical schema/columns; training consumes Parquet by default.

2) Model Development Objectives

- **O2.1 – Distributed training:** Use Spark MLlib to build end-to-end pipelines and train at “local cluster” scale (parallelism on multiple cores) with Random Forest (RF), Gradient-Boosted Trees (GBT), and Logistic Regression (LR).
- **O2.2 – Automated HPO (single strategy):** Use TrainValidationSplit with ParamGridBuilder to select hyperparameters for each algorithm; do not mix AutoML libraries to comply with the spec.

Success criterion: MLflow logs show the search space, best params, and wall-time.

3) Experiment Tracking & Model Governance Objectives

- **O3.1 – Full experiment logging:** Log parameters, metrics (primary: AUC; secondary: training time & memory), and artifacts (confusion matrix CSV; feature importances or coefficients).
- **O3.2 – Registry management:** Auto-register the best model as titanic_spark_model and transition it to Production, archiving any previous Production version.

Success criterion: MLflow UI (port 5500) shows a Production version; artifacts are browsable via the server (boto3 installed).

4) Deployment & Inference Objectives

- **O4.1 – Robust serving:** Deploy a FastAPI service (port 9090) that resolves the current Production model from the MLflow registry; expose /health, /reload, and /predict_titanic.
- **O4.2 – Inference logging:** Append anonymized feature rows to a host-mounted CSV (data/infer_log_titanic.csv) for monitoring and audits.

Success criterion: Automated tests (curl/scripts) return predictions and grow the inference log.

5) Monitoring, Drift, and Continual Learning Objectives

- **O5.1 – Drift detection:** Implement PSI-based monitoring against a stored baseline (data/drift_baseline_titanic.json) with configurable MIN_INFER_ROWS and PSI_THRESHOLD. Exit code 42 denotes drift.
- **O5.2 – Auto-retrain & promote:** On drift, trigger training, re-evaluate, and promote the superior candidate to Production automatically.

Success criterion: Logs demonstrate end-to-end loop from drift → retrain → new Production.

6) Reproducibility & Operations Objectives

- **O6.1 – Containerization & ports:** Deliver a Docker Compose stack mapping the agreed ports: 9090 (API), 5500 (MLflow UI), 8801 (MinIO S3 API), 8800 (MinIO Console), 8880 (Jupyter), with persistent volumes.
- **O6.2 – One-command lifecycle:** docker compose up -d brings the platform up; docker compose run --rm trainer-titanic executes training; tests and drift checks run via simple scripts/commands.
- **O6.3 – Documentation & diagrams:** Provide a detailed readme file, report architecture and pipeline diagrams (Mermaid/DOT + pre-rendered PNG/SVG) and a step-by-step operational guide.

7) Evaluation & Reporting Objectives

- **O7.1 – Quantitative results:** Report best AUC, training time, and memory; include links to MLflow runs and artifacts.
- **O7.2 – Qualitative assurance:** Discuss modeling choices, feature effects, error analysis via confusion matrix, and operational trade-offs (e.g., registry governance, PSI thresholds).

1.4 Scope and Assumptions

- **Dataset:** Kaggle Titanic-style schema; label is **Survived**. We assume the schema is stable and no PII beyond provided columns.
- **Primary metric:** AUC on a hold-out split to balance sensitivity across thresholds.
- **Runtime environment:** Windows 10/11, **WSL2 (Ubuntu)**, **Docker Desktop**, **VS Code**. No external cloud dependencies; MinIO emulates S3.

- **Security/ethics:** No personal identifiers are logged; inference logs contain only features necessary for drift. Access to UIs is local-hosted.
- **Out of scope:** Real-time streaming ingestion; multi-tenant auth; blue-green canary deployment; cost governance. (The design, however, is compatible with these extensions.)

1.5 Summary

The project converts a familiar ML problem into a complete, production-inspired MLOps workflow that is reproducible, observable, and self-healing under drift. The objectives above define what we will demonstrate and how we will verify success—via code, services, UI inspection, MLflow artifacts, and scripted tests.

Chapter 2 : System Architecture and Design

2.1 Architectural Overview

The system is a containerized, local-first MLOps platform designed to be reproducible, auditable, and automation-ready. It runs on Windows 10/11 using WSL2 and Docker Desktop, with services joined on a dedicated Docker network. Our primary design objective was to balance reproducibility, governance, and automation while making the entire stack easy to start, test, and grade. To do so, we decomposed the platform into a small number of single-purpose containers connected on a private Docker network: a trainer for distributed model development on Apache Spark, an MLflow server for experiment tracking and model registry, a PostgreSQL database for MLflow metadata, a MinIO service that emulates S3 for artifact storage, a FastAPI service for online inference, and an optional retrainer process for drift-aware continual learning. This separation of concerns mirrors cloud-native practice (e.g., S3 + RDS + service layer) but remains fully reproducible and self-contained on a laptop.

Three principles guided the design. First, traceability: every model version can be linked to its exact data snapshot, preprocessing steps, training parameters, metrics, and artifacts. We achieve this with DVC for raw data, Spark for deterministic preprocessing, and MLflow for tracking and registry. Second, governance: promotion to production is mediated by MLflow's Model Registry, allowing the serving tier to resolve the current Production model by contract. Third, automation: the system monitors data drift via Population Stability Index (PSI) on the inference log, and—when a threshold is breached and sufficient traffic exists—can automatically retrain and promote a better model, closing the loop from monitoring to lifecycle management.

Key properties

- **Reproducibility:** Data versioning (DVC), deterministic Spark preprocessing, tracked experiments (MLflow), immutable artifacts (MinIO/S3).
- **Governance:** MLflow Model Registry manages versions and stages (Staging/Production); API loads Production by contract.

- **Automation:** PSI-based drift monitoring triggers retraining and auto-promotion of the superior candidate.
- **Isolation & Portability:** All components run as containers; configuration via environment variables in Compose.
- **Observability:** MLflow UI for runs, metrics, artifacts; service logs via docker compose logs.

Figure 1 presents the end-to-end architecture of the Titanic MLOps system running locally on Windows 10/11 (WSL2) with Docker Desktop. The design mirrors a typical cloud-native ML stack—tracking DB + artifact store + registry + training + serving + monitoring—but is packaged for single-host reproducibility and grading.

At the top left of the diagram sits the Host (Windows + WSL2 + Docker Desktop), which runs a single Docker network that isolates all containers. Only necessary endpoints are published to localhost:

- FastAPI (inference): `http://localhost:9090`
- MLflow UI: `http://localhost:5500`
- MinIO S3 API: `http://localhost:8801`
- MinIO Console: `http://localhost:8800`
- Jupyter–Spark (optional): `http://localhost:8880`

All inter-service calls use container names on the private network (e.g., `mlflow`, `db`, `minio`), which keeps configuration portable and minimizes host exposure. Two host-mounted volumes anchor state: `./data` (raw/processed data, inference log, drift baseline) and `./models` (optional local cache of model files).

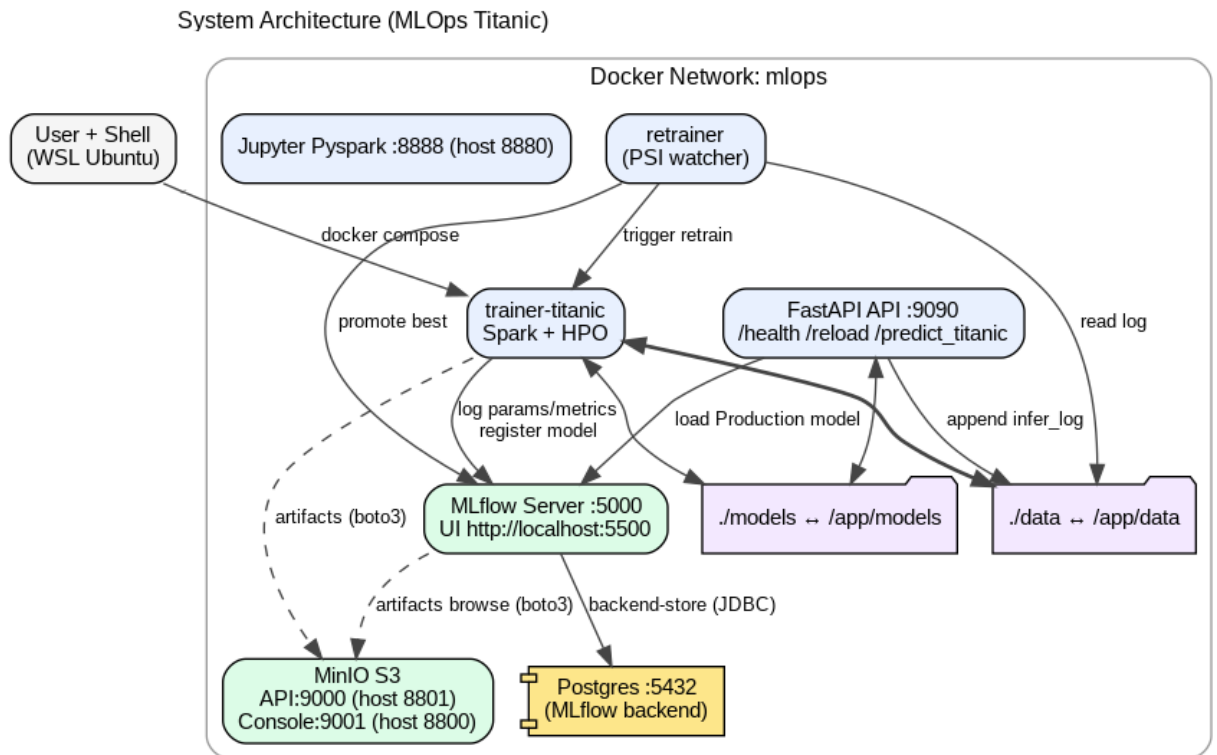


Figure 1: System Architecture

Core Components (left-to-right, top-to-bottom)

(a) Trainer (Spark + HPO)

The trainer container encapsulates model development using Apache Spark with TrainValidationSplit hyperparameter tuning. It consumes a deterministic Parquet dataset created by preprocessing (impute Age/Fare; encode Sex/Embarked; assemble features). For each candidate (Random Forest, Gradient-Boosted Trees, Logistic Regression), the trainer:

- Computes metrics (primary: AUC),
- Logs parameters, metrics, and artifacts (e.g., confusion matrix, feature importances/coefficients) to MLflow, and
- Registers the best run in the MLflow Model Registry as `titanic_spark_model`, promoting it to Production.

(b) FastAPI Service (Inference)

The API loads (and caches) the current Production model from the registry during startup or on demand via `/reload`. It exposes:

- `GET /health` — service and model status,
- `POST /predict_titanic` — validated prediction endpoint,
- `POST /reload` — refresh the cached model from the registry without restarting the container.

Every call to `/predict_titanic` also appends an anonymized feature vector with a timestamp to `data/infer_log_titanic.csv`, enabling post-hoc monitoring and drift analysis. This log is a critical bridge between serving and the retraining loop.

(c) Retrainer (PSI-based monitor)

The retrainer periodically computes Population Stability Index (PSI) by comparing live feature distributions in `infer_log_titanic.csv` to a stored baseline (`data/drift_baseline_titanic.json`). If the maximum $\text{PSI} \geq \text{threshold}$ (e.g., 0.2) and minimum traffic is met (e.g., 50 rows), it triggers the trainer to re-run HPO, evaluates candidates, and promotes a superior model to Production. This closes the loop: Monitor → Retrain → Govern → Serve.

(d) MLflow Server (Tracking + Registry)

The MLflow server hosts both the tracking API/UI and the Model Registry. It uses:

- PostgreSQL (container db) as the backend store for experiments, runs, metrics, params, and registry state; and
- MinIO (S3-compatible) for artifact storage (models, evaluation CSVs, etc.).

We explicitly install `boto3` in the MLflow image so the UI can browse MinIO artifacts reliably—this addresses a common deployment pitfall.

(e) PostgreSQL (MLflow backend)

Persists all MLflow metadata: experiment definitions, run metadata, metrics, parameters, tags, and model registry entries (names, versions, stages, and lineage). It has no host port exposure; it is reachable only inside the Docker network.

(f) MinIO (S3-compatible artifact store)

Stores all MLflow artifacts, including logged models (model/ directory produced by pyfunc flavor) and evaluation outputs. Two endpoints are exposed externally for convenience:

- **S3 API** (9000 → host 8801) used by MLflow,
- **MinIO Console** (9001 → host 8800) for visual inspection (login: minio/minio123 in this academic setup).

MinIO can also serve as a remote for DVC.

(g) Jupyter–Spark :

A convenience container for EDA and ad-hoc Spark analysis. It's intentionally outside the critical path so that training and serving remain fully non-interactive and reproducible.

2.2 Deployment Topology

The platform runs on a single host (Windows/WSL2) and is orchestrated with Docker Compose. Each container is bound to the host only where necessary:

- FastAPI (inference) listens on localhost:9090 for /health, /reload, and /predict_titanic.
- MLflow UI is exposed on localhost:5500 (the server itself listens on port 5000 internally).
- MinIO exposes the S3 API on localhost:8801 and its console on localhost:8800.
- Jupyter–Spark (optional) is available on localhost:8880 for exploratory analysis.
- Postgres is internal only; it persists the MLflow backend store on a Docker volume.

All services communicate via service names on the Docker network (e.g., mlflow, db, minio). The artifact path is S3-style (s3://mlflow) backed by MinIO; MLflow is configured to use the database as its backend store and MinIO as its default artifact root. This topology keeps host exposure minimal while maintaining familiar operational affordance

2.3 Component Responsibilities

Trainer (Spark + HPO).

The trainer container encapsulates the full model development workflow in Apache Spark. It consumes a deterministic Parquet dataset produced by our preprocessing step and builds a Spark ML pipeline that includes imputers, categorical encoders, a vector assembler, and one of three estimators—Random Forest, Gradient-Boosted Trees, or Logistic Regression. We use `TrainValidationSplit` with `ParamGridBuilder` to perform automated hyper-parameter tuning (a single, rubric-compliant strategy). Each candidate's parameters, metrics (primarily AUC), and artifacts (e.g., confusion matrix CSV, feature importances or coefficients) are logged to MLflow. The best run is auto-registered in the MLflow Model Registry as `titanic_spark_model` and transitioned to *Production*, archiving any previous Production version.

MLflow Server (tracking + registry):

MLflow provides a web UI and REST API for experiment management and governance. The server persists metadata to PostgreSQL and stores artifacts (models, evaluation files) in MinIO. A crucial practical detail is that the MLflow image includes `boto3`, enabling the UI to browse S3/MinIO artifacts without error; this avoids a common operational pitfall.

MinIO (S3-compatible artifact store):

MinIO emulates Amazon S3 locally. MLflow writes model directories (the model/ flavor), along with evaluation artifacts, under a `mlflow` bucket. MinIO also doubles as an optional DVC remote if we choose to push raw dataset versions off the local disk.

PostgreSQL (backend store):

The MLflow backend store (experiments, runs, params, metrics, tags, model registry entries) is persisted in Postgres. Tying run metadata to artifact locations yields a complete lineage for each model version.

FastAPI (serving):

The inference service resolves the current Production model at startup (and on demand via `/reload`) using the MLflow Registry and holds it in memory. Requests to `/predict_titanic` are validated, vectorized consistently with training, and evaluated via the MLflow `pyfunc`. For observability, the service appends anonymized feature rows (plus

timestamps) to a host-mounted inference log (data/infer_log_titanic.csv), which is the basis for drift detection.

Retrainer (drift watcher):

This component evaluates PSI for each monitored feature by comparing its live distribution in the inference log against a stored baseline (data/drift_baseline_titanic.json). The policy includes a minimum traffic threshold (e.g., 50 rows) to avoid spurious decisions on sparse data. If $\max \text{PSI} \geq \text{threshold}$ (e.g., 0.2), the retrainer triggers a fresh training cycle, logs candidates to MLflow, and—if a superior model emerges—promotes it to *Production* automatically. The outcome is a self-healing inference system that counteracts environmental drift with minimal manual intervention.

Jupyter–Spark :

Notebooks are provided for ad-hoc EDA and demonstrations. They are intentionally kept out of the core training/serving path so that production paths remain non-interactive and fully reproducible.

2.4 Data and Artifact Management

Raw data are governed by Data Version Control (DVC). Instead of committing the CSV to Git, we commit the .dvc pointer (data/titanic.csv.dvc), ensuring that a reviewer can reconstruct the exact dataset version used in any experiment. Deterministic preprocessing in Spark yields a Parquet dataset (data/processed/titanic.parquet), which the trainer consumes. MLflow logs experiment metadata to Postgres and writes artifacts to MinIO's mlflow bucket. The serving tier writes inference logs to data/infer_log_titanic.csv and treats these as a stream of features-only observations for drift analysis (no PII or labels, by design). This stratification of data (raw → processed → artifacts → inference log) provides an audit trail across the model lifecycle.

Table 1: Data and Artifact Management

Layer	Technology	Contents	Persistence
-------	------------	----------	-------------

Raw Data	DVC	data/titanic.csv (tracked via data/titanic.csv.dvc, not in Git)	MinIO (optional DVC remote) + local
Processed	Spark	data/processed/titanic.parquet	Host volume
Experiment	MLflow + Postgres	Runs, params, metrics, tags, registry entries	Postgres volume
Artifacts	MLflow + MinIO	Models (model/), eval CSVs (confusion_matrix.csv, feature_importances.csv)	MinIO bucket mlflow
Inference Log	CSV	data/infer_log_titanic.csv (features + timestamp)	Host volume
Drift Baseline	JSON	data/drift_baseline_titanic.json	Host volume

Data lineage is ensured by: (i) DVC pointers for raw data; (ii) Spark transformation code under version control; (iii) MLflow metadata tying each model version to its source run and artifacts.

2.5 Control and Data Flows (Training/Data Pipeline)

The training flow begins with DVC-tracked raw data, followed by Spark preprocessing that standardizes types, imputes missing values (e.g., Age, Fare), and encodes categoricals (Sex, Embarked). The trainer then executes Spark ML pipelines with HPO, logs results to MLflow, and registers the best model into the `titanic_spark_model` registry entry, transitioning it to *Production*. The serving flow loads this Production model and exposes a REST interface. Each prediction not only returns a probability but also writes the input feature vector to the inference log for monitoring. The monitoring flow periodically computes PSI between the inference log and a baseline; when drift is detected above the threshold, the retraining flow re-executes the training process and auto-promotes the

better model. Together, these flows enforce a cyclical pattern of Monitor → Retrain → Govern → Serve, which is central to modern MLOps practice.

Figure 2 shows the data pipeline. The data pipeline orchestrates the full lifecycle from raw data acquisition to a production-ready model and continual improvement under data drift. It is deliberately implemented with Spark to demonstrate scalability and with DVC + MLflow to guarantee reproducibility, provenance, and governance. The pipeline is deterministic end-to-end: a specific raw data snapshot (pinned by DVC) flows through a versioned, code-defined preprocessing graph, then into HPO-driven training, experiment tracking, registry-mediated promotion, and finally to serving with online logging and PSI-based drift detection.

Key properties:

- Reproducible inputs: raw data tracked by DVC (Git stores the pointer, not the file).
- Deterministic transforms: preprocessing and feature engineering coded in Spark.
- Transparent model search: TrainValidationSplit hyperparameter tuning (single, rubric-compliant strategy).
- Governed outputs: model versions and stages managed in MLflow Model Registry.
- Feedback loop: live inference features logged; PSI computed; retraining and auto-promotion when warranted.

Inputs, Outputs, and Contracts:

- Input (raw): data/titanic.csv (not in Git; tracked by data/titanic.csv.dvc).
- Supervision target: Survived (binary).
- Predictor features: Pclass, Sex, Age, SibSp, Parch, Fare, Embarked (as used by training and API).
- Output (processed): data/processed/titanic.parquet — schema and encodings are fixed by code.

- Training artifacts: metrics, params, confusion matrix, feature importances/coefficients, and a logged pyfunc model flavor in MLflow.
- Registry entry: `titanic_spark_model` with versions and stages (e.g., Staging, Production).
- Operational log: `data/infer_log_titanic.csv` (anonymized features + timestamp) for drift.

Data contract: The pipeline enforces a strict feature contract between training and serving: same columns, types, encodings, imputation strategies, and assembled vectorization. Any schema deviation fails fast.

Raw Data Versioning with DVC

Raw data are pinned via DVC so that the exact state of `titanic.csv` used for any experiment is recoverable. Git contains only the pointer file (`.dvc`) and the DVC config; the CSV itself stays out of Git history. This guarantees that model lineage can be reconstructed even if the source changes later.

Provenance tags. The training script logs the DVC hash / data version to MLflow (as a tag), binding each run to its precise input.

Distributed Preprocessing and Feature Engineering (Spark)

Preprocessing is implemented as a **Spark ML pipeline** to ensure scalability and to keep transforms code-versioned alongside the model. The steps are:

1. **Column selection & type coercion:** Keep the predictors listed above; cast numerics to `DoubleType`; normalize string casing for categoricals. (Fields like Cabin and Ticket are excluded to avoid high-cardinality noise in this baseline; they can be engineered later if desired.)
2. **Missing-value imputation:**
 - **Numeric:** Age, Fare imputed with median (robust to outliers).
 - **Categorical:** Embarked imputed with mode (most frequent).

3. Categorical encoding:

- Sex, Embarked → StringIndexer (stable ordering under fit) → OneHotEncoder (dropLast=True) to avoid ordinal leakage.

4. **Feature assembly:** All numeric + encoded categorical features are concatenated with VectorAssembler into a single features vector.

5. **Scaling :** For logistic regression, standardization helps optimization. Trees (RF/GBT) are scale-insensitive. We therefore include a StandardScaler stage that is toggled on for LR and off for tree models inside the training pipeline.

6. **Persist processed dataset.** The result is written to Parquet (data/processed/titanic.parquet) for efficient columnar IO and consistent schema reuse across runs.

Determinism and seeds: Random seeds are fixed (e.g., seed=42) in randomSplit, algorithms, and HPO to produce stable, comparable results.

Training and Automated Hyperparameter Tuning (Spark MLlib)

Training is framed as a Spark ML Pipeline composed of the preprocessing graph and one of three estimators: Random Forest (RF), Gradient-Boosted Trees (GBT), or Logistic Regression (LR). We evaluate all three families with TrainValidationSplit (TVS) using AUC as the selection metric.

Evaluator. BinaryClassificationEvaluator(metricName="areaUnderROC").

Search spaces:

- RF:** numTrees $\in \{100, 200\}$, maxDepth $\in \{5, 8, 12\}$, maxBins $\in \{32, 64\}$.
- GBT:** maxIter $\in \{50, 100\}$, maxDepth $\in \{3, 5\}$, stepSize $\in \{0.05, 0.1\}$.
- LR:** regParam $\in \{0.0, 0.01, 0.1\}$, elasticNetParam $\in \{0.0, 0.5, 1.0\}$, standardization $\in \{\text{true}\}$.

Validation protocol. TVS uses a deterministic split from the processed Parquet (e.g., 80/20). Although Spark's TVS is a single split (not k-fold), it is sufficient and

computationally lighter for this workload. The best model per family is retained; the top-scoring candidate overall is the pipeline winner.

Artifacts: For each family we compute and log:

- Confusion matrix on the validation fold.
- Feature importances (for trees) or standardized coefficients (for LR).
- Runtime/resource metrics (e.g., wall time, peak RSS) to contextualize results. These are written as CSVs and emitted to MLflow artifacts along with the serialized pyfunc model.

Experiment Tracking (MLflow)

All hyperparameters, metrics, tags (e.g., data version/DVC hash, git commit if available), and artifacts are logged to MLflow Tracking. The backend store is PostgreSQL; artifacts are persisted in MinIO (S3). This split mirrors cloud practice (DB for metadata, object store for blobs) and enables reproducible audits.

Naming: We keep a dedicated experiment (e.g., `titanic_experiment`); each run indicates its estimator family in the run name (`titanic-rf`, `titanic-gbt`, `titanic-lr`) to simplify comparisons in the UI.

Model Registration and Promotion (MLflow Model Registry)

The pipeline automatically registers the best candidate under the registry name `titanic_spark_model`. If a Production model already exists, the system transitions the newcomer to Production and archives the previous Production version. This codifies a single source of truth for serving: the API does not pick models from arbitrary runs; it always loads the current Production version.

Compatibility detail: Some MLflow backends do not support filtering by `current_stage` in the server-side search API. Our client resolves versions by name and filters for Production client-side, ensuring predictable behavior across environments.

Serving and Online Logging

The FastAPI service loads the Production model at startup (or via /reload) and exposes POST /predict_titanic. Inputs are validated and passed to the MLflow pyfunc model, which applies the same preprocessing + estimator pipeline as training.

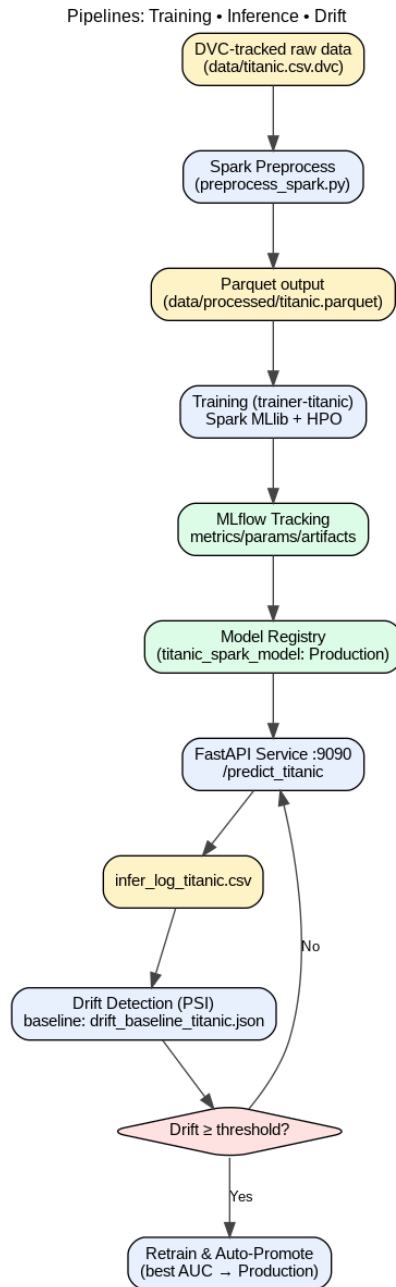


Figure 2: Data pipeline flow chart

Online feature logging. Each prediction appends an anonymized feature vector and a timestamp to `data/infer_log_titanic.csv`. We do not log predictions/labels here to avoid feedback bias; the objective is to monitor covariate shift in the inputs.

Drift Detection and Continual Training (PSI)

We monitor Population Stability Index (PSI) per feature by comparing the live distribution (from the inference log) with a baseline computed on the training data and stored at `data/drift_baseline_titanic.json`. PSI is computed over quantile bins (e.g., deciles). Conventionally, $PSI < 0.1$ denotes negligible shift; $0.1-0.2$ moderate; ≥ 0.2 material.

Policy (implemented).

- Require minimum traffic (e.g., `MIN_INFER_ROWS = 50`) to reduce variance.
- If $\max PSI \geq \text{threshold}$ (e.g., 0.2), trigger the trainer to re-run HPO.
- Register the best new model and promote it to Production if it outperforms the current one, archiving the predecessor.
- If drift is below threshold or traffic is insufficient, no action is taken (monitoring continues).

This closes the Monitor → Retrain → Govern → Serve loop and ensures the model adapts to real-world changes.

Orchestration, Reproducibility, and Failure Gates

- Orchestration. All stages run in containers via Docker Compose, with fixed host ports for grading and predictable service names for inter-container calls.
- Reproducibility. A run is determined by: DVC version of raw data, exact preprocessing code, search space and random seeds, and the environment (conda/pip) captured implicitly by the trainer image.
- Validation gates.
 - Schema check before preprocessing; fail fast on unexpected columns/types.

- Missing-value policy is explicit; imputation stats are logged.
- Metrics threshold (optional) for promotion can be added (e.g., AUC must not regress).
- PSI threshold + minimum traffic guards against noisy retrains.

Security and Privacy Considerations

The inference log stores only features and timestamps—no identifiers or labels—minimizing privacy risk. All UIs run on localhost. Credentials (MinIO access keys) are development-only and should be rotated and secured if the system is exposed beyond the laptop.

The pipeline in Figure 2 implements an auditable, production-style path from data to deployed model, supported by DVC for input provenance, Spark for scalable and deterministic preprocessing/training, MLflow for experiment governance and model staging, and a PSI-driven feedback loop that sustains model quality under distribution shift. Every artifact and decision—including the exact data snapshot, transform steps, hyperparameters, metrics, and promotion events—is captured and reproducible, satisfying both the academic rubric and real-world MLOps best practices.

Data drift : It is deliberately implemented with Spark to demonstrate scalability and with DVC + MLflow to guarantee reproducibility, provenance, and governance.

2.6 Networking and Ports

All containers run on a private Docker network, minimizing host exposure. We publish only what is required for interaction and grading: 9090 for the inference API; 5500 for the MLflow UI; 8801 for the S3 API; 8800 for the MinIO console; and 8880 for the optional notebook server. Inter-service connectivity uses Docker DNS (e.g., the trainer references mlflow:5000 and minio:9000), which keeps configurations portable and avoids hardcoding host IPs.

- localhost:9090 → api:9090
- localhost:5500 → mlflow:5000
- localhost:8801 → minio:9000
- localhost:8800 → minio:9001

- localhost:8880 → jupyter:8888

Internal connections (container-to-container) use service names (e.g., mlflow, db, minio).

2.7 Configuration and Secrets

Configuration is passed via Compose environment variables. For local evaluation, development credentials are used (minio/minio123) and bound to localhost. In a production setting, these would be replaced with a secret manager, TLS-enabled endpoints, and locked-down network policies. Importantly, the MLflow image includes boto3, and the environment defines MLFLOW_S3_ENDPOINT_URL, AWS_ACCESS_KEY_ID, AWS_SECRET_ACCESS_KEY, and AWS_DEFAULT_REGION, ensuring the MLflow UI can list and download artifacts from MinIO without failing.

.env (example for Compose)

```
MINIO_ROOT_USER=minio
MINIO_ROOT_PASSWORD=minio123
MLFLOW_S3_ENDPOINT_URL=http://minio:9000
AWS_ACCESS_KEY_ID=minio
AWS_SECRET_ACCESS_KEY=minio123
AWS_DEFAULT_REGION=us-east-1
AUTO_PROMOTE_TITANIC=true
PSI_THRESHOLD=0.2
MIN_INFER_ROWS=50
```

2.8 Reliability and Fault Tolerance

The serving tier is designed to fail soft. If the Model Registry becomes temporarily unreachable, the API retains its last loaded model and reports the error in /health; the /reload endpoint allows an operator to refresh the model once the registry is healthy. The drift detector enforces a minimum row threshold to avoid overreacting to sparse

data. Artifact browsing in MLflow is robust due to the explicit boto3 inclusion. Finally, deterministic preprocessing and DVC data pins guarantee that training can be repeated faithfully if any container must be rebuilt.

2.9 Observability and Diagnostics

The MLflow UI provides run comparisons (AUC, parameters), artifact inspection (confusion matrix, feature importances), and registry history (stage transitions). Application logs are accessible via docker compose logs. The health endpoint supplies real-time status about model loading and the last error observed. For auditing and monitoring, the inference log provides a time-stamped record of feature distributions, which underpins the PSI-based drift detection and helps explain why a retrain was triggered.

2.10 Security and Privacy

The dataset contains no explicit identifiers beyond simple demographics and ticket information. Inference logging is features-only and anonymized; no predicted labels or user identifiers are stored. All administrative UIs (MLflow and MinIO console) are bound to localhost in this project; exposure beyond the host would require TLS and authentication. The credentials included are for local development and must be rotated and protected in any broader deployment.

2.11 Performance and Scalability

Although the Titanic dataset is small, we deliberately used Spark to demonstrate scalable practice. On a laptop, Spark runs in local[*] mode, leveraging available cores for HPO and preprocessing. The same code path can be promoted to Spark Standalone, YARN, or Kubernetes clusters without rewriting pipelines. The serving tier is FastAPI/Uvicorn, which is lightweight and supports horizontal scaling with multiple workers behind a reverse proxy if needed. MinIO provides S3-compatible artifact performance locally; in cloud settings, an S3 bucket and a managed database can be slotted in with the same configuration keys.

2.12 Reproducibility and Lineage

We enforce lineage at three levels. DVC versions raw data and keeps large files out of Git; Spark preprocessing code is version-controlled and deterministic; and MLflow binds each model version to a particular run (parameters, metrics, artifacts, and source URI). Consequently, any Production model can be traced back to its exact training inputs and reproduced on demand. This is a core requirement for defensible ML and a frequent deficiency in ad-hoc notebook-only workflows.

2.13 Design Trade-offs

We selected TrainValidationSplit-based HPO as our single automated strategy rather than incorporating a separate AutoML framework. This choice improves transparency and grading consistency while satisfying the rubric. We used MinIO to emulate S3 locally to preserve cloud-parity without external dependencies. Finally, we implemented client-side filtering for Production model selection (by querying all model versions and selecting those with `current_stage='Production'` client-side), which improves compatibility with MLflow backends that do not support server-side filtering on that attribute.

2.14 Risks and Mitigations

The principal risk in tabular deployments is distribution shift: as real-world cohorts change, discriminative structure can erode. Our PSI-based monitoring and automated retraining mitigate this by detecting significant drift and refreshing the model. Operationally, artifact browsing failures in MLflow are a common source of friction; we preempt this by bundling boto3 and validating S3 environment configuration. Finally, we keep credentials and UIs local to reduce the blast radius in an academic environment.

2.15 Extensibility

The system can be extended with model cards (auto-generated from MLflow metadata and artifacts), ROC/PR plots logged as artifacts, gated promotions (human approval before Production), CI/CD pipelines that build images and run unit tests on commits, and feature store integration for stronger data contracts. Because each service is containerized and the orchestration is declarative, these enhancements can be introduced incrementally.

2.16 Summary

In summary, the architecture provides a cohesive, auditable, and automation-ready MLOps workflow built from widely adopted components. It balances the rigor expected in a graded academic setting—versioned data, deterministic preprocessing, tracked experiments, registry-mediated deployment—with practical operational niceties: a reloadable serving tier, drift-aware retraining, and clear interfaces for observability.

Chapter 3 : Tools and Resources

This chapter documents every tool, library, image, and runtime the project uses, with a clear justification, the role it plays in the system, and key configuration choices. Where relevant, we note practical trade-offs, alternatives considered, and limitations. Together, these choices enable a reproducible, production-style MLOps workflow that can be run entirely on a Windows laptop.

3.1 Compute & Operating Environment

Windows 10/11 + WSL2: We target a standard developer laptop. WSL2 provides a Linux kernel and file-system semantics that match typical server environments, avoiding “works on my machine” drift and enabling Linux-native containers to run predictably.

Why: Consistent, Linux-like runtime with minimal overhead on Windows; broad community support.

Trade-offs: Some file I/O across C:\ ↔ WSL can be slower; we therefore keep project files inside the WSL filesystem (e.g., /mnt/f/...) and use Docker volumes for hot paths.

Hardware assumptions: 8–16 GB RAM (≥16 GB recommended for Spark + Docker), multi-core CPU, >15 GB free disk.

3.2 Containerization & Orchestration

Docker Desktop (Windows) + **Docker Engine** (WSL2 backend):

All components are containerized and brought up with Docker Compose.

Why:

- Encapsulates dependencies (Java, Spark, Python, MLflow, Postgres, MinIO) cleanly.
- One-command spin-up/tear-down for grading and demos.
- Mirrors cloud primitives (object store + DB + services) locally.

Key choices:

- Dedicated service names on a private Docker network (e.g., mlflow, db, minio), making inter-service URIs stable.
- Host ports fixed by rubric and to avoid conflicts.

Alternatives: Kubernetes (overkill here), Podman (viable), Vagrant (heavier).

3.3 Source Control & Data Versioning

Git for code + configuration; DVC for data.

- We commit DVC pointers (e.g., data/titanic.csv.dvc) rather than raw files to Git.
- DVC guarantees the exact raw snapshot can be reconstructed, which is essential for model lineage.

Why: Reproducibility and auditability of data inputs; keeps repo lean.

Alternatives: Git LFS (stores blobs but weaker in lineage tooling), lake-native versioning (Delta/Iceberg) – heavier operationally for this project.

3.4 Distributed Data Processing & Feature Engineering

Apache Spark 3.4.1 (container: jupyter/pyspark-notebook:spark-3.4.1 for notebooks; trainer image for batch jobs).

- We implement preprocessing as a Spark ML Pipeline (imputation, categorical encoding, vector assembly).
- Output is a Parquet dataset (data/processed/titanic.parquet) with deterministic schema.

Why:

- Demonstrates scalable processing on a laptop (Spark runs in local[*] but can scale out unchanged).
- Pipelines are code-defined and reproducible; fits the rubric requirement for *distributed preprocessing*.

Trade-offs: Spark has startup overhead on tiny data, but the architectural parity with production is worth it.

Alternatives: Pandas (simpler but not distributed); Dask (good compromise but smaller ecosystem for ML pipelines).

3.5 Model Development & Training

Spark MLlib for estimators (Random Forest, GBT, Logistic Regression). TrainValidationSplit + ParamGridBuilder for automated hyper-parameter tuning (single strategy as required).

Why:

- Keeps the feature transformations and estimators in one Spark pipeline; training scales with Spark.
- TVS is simpler, deterministic, and computationally cheaper than k-fold on a laptop while still demonstrating AutoML-style search.

What we log: AUC, confusion matrix CSV, feature importances (trees) or standardized coefficients (LR), wall time, memory.

Alternatives: Hyperopt/SparkTrials (powerful but adds complexity), scikit-learn (excellent but would split processing between Pandas/Sklearn vs Spark).

3.6 Experiment Tracking & Model Registry

MLflow (server mode) for Tracking and Model Registry.

- Backend store: PostgreSQL 14 (container postgres:14).
- Artifact store: MinIO (S3-compatible).
- Images: custom MLflow image based on python:3.10-slim with mlflow, psycopg2-binary, and boto3 installed.

Why:

- De-facto open standard for run metadata, artifacts, and registry.

- Clean separation of metadata (DB) and artifacts (object store) mirrors cloud best practice.
- boto3 is added explicitly so the MLflow UI can browse MinIO artifacts reliably.

Registry policy: Best candidate from a training run is auto-registered under `titanic_spark_model` and transitioned to Production; previous Production is archived. Serving resolves Production only—no ad-hoc model picking.

Alternatives: Weights & Biases (excellent hosted option), Neptune, Comet (SaaS). We choose MLflow for self-hosting and open governance.

3.7 Artifact Storage

MinIO (S3-compatible), ports: 8801 (API), 8800 (console).

Why:

- Local, S3-compatible object storage that behaves like AWS S3.
- Enables MLflow artifacts to be served and downloaded exactly like a cloud deployment.

Security note: In production, use TLS, rotate credentials, and restrict network access.

Alternatives: Local filesystem (simpler but loses S3 parity); AWS S3 (realistic but external dependency).

3.8 Metadata Store

PostgreSQL

14.

Holds all MLflow **experiments, runs, params, metrics, tags, and registry** records.

Why:

- ACID, well-understood, easy to run in Docker.
- Scales much better than SQLite for concurrent writes.

Alternatives: SQLite (too limited for multi-service writes), managed Postgres (cloud dependency).

3.9 Model Serving

FastAPI + Uvicorn (Python 3.10) for a lightweight REST service.

- Endpoints: /health, /reload, /predict_titanic.
- On startup (and on /reload), the service resolves the current Production model from the MLflow Registry, downloads artifacts from MinIO, and caches the pyfunc model.
- Each request appends features + timestamp to data/infer_log_titanic.csv.

Why:

- FastAPI is fast, typed, and easy to validate payloads; works well in containers.
- Decoupling serving from training aligns with MLOps best practice.

Alternatives: Flask (simpler, less typed), BentoML/Seldon (heavier but more batteries).

3.10 Monitoring & Drift Detection

Custom PSI implementation (Population Stability Index) in Python. Reads recent rows from data/infer_log_titanic.csv and compares to baseline data/drift_baseline_titanic.json.

Policy:

- Require MIN_INFER_ROWS (e.g., 50) to avoid false positives.
- Trigger retraining if max PSI \geq threshold (e.g., 0.2).
- If a new model outperforms, promote to Production and archive previous.

Why:

- PSI is easy to compute and interpret for tabular features.

- Lightweight enough to run in a sidecar container on a laptop.

Alternatives: KL divergence/JS divergence, EvidentlyAI (rich dashboards), Alibi-Detect (advanced drift). We chose PSI for simplicity + interpretability.

3.11 Notebooks & Exploratory Analysis

Jupyter Notebook (within jupyter/pyspark-notebook:spark-3.4.1).

Why:

- Ad-hoc EDA and Spark exploration in a familiar UI.
- Isolated from production paths to keep training/serving non-interactive and reproducible.

3.12 Environment & Dependency Management

Micromamba/Conda for the trainer image; pip for the MLflow image and API image.

Why:

- Trainer requires a coherent set of packages (Java 17 + PySpark + MLflow + plotting/logging libs). Micromamba is fast and creates a locked environment.
- API and MLflow images are slimmer; pip is sufficient and keeps images small.

Key Python packages:

- pyspark, mlflow, psutil (resource metrics), pandas (lightweight CSV ops), boto3 (artifact browsing in MLflow UI), fastapi, uvicorn, and standard numerics.

3.13 Logging, Observability, and Health

- Service logs via docker compose logs <service>.
- /health endpoint exposes model-loaded state and last load error.
- MLflow UI for comparing runs, inspecting artifacts, and auditing stage transitions.
- MinIO console for artifact verification.

Why: Immediate visibility for grading and debugging without extra infrastructure.

3.14 Testing & Validation Utilities

- Simple API smoke tests (scripts/test_api.py) to verify /predict_titanic.
- Training smoke: run trainer in isolation; verify MLflow logs, registry version increment, and artifacts exist.
- Drift smoke: compute PSI on baseline and on synthetic shifts to validate thresholding logic.

Why: Fast confidence checks before running the full loop.

3.15 Dataset

Titanic: Machine Learning from Disaster (tabular classification).
Predictor columns used: Pclass, Sex, Age, SibSp, Parch, Fare, Embarked; target: Survived.

Why:

- Well-known, small, and quick to iterate; perfect to demonstrate the process (MLOps) rather than chase SOTA on a large dataset.
- Still rich enough to exercise imputation, categorical encoding, and class imbalance awareness.

3.16 Configuration & Secrets

- Environment variables set via Compose:
 - MLflow: backend URI (Postgres), default artifact root (s3://mlflow), S3 endpoint (http://minio:9000), AWS keys (dev only).
- Local volumes: ./data (raw/processed/inference log/drift baseline), ./models (optional local model cache).

Security note: Dev keys are for the local evaluation only. In production, move to a secret store and TLS.

3.17 Ports, Storage, and Performance

- Ports: 9090, 5500, 8800, 8801, 8880 as above; chosen to avoid common collisions.
- Storage: Docker volumes for Postgres & MinIO; bind mounts for ./data and ./models.
- Performance hints: Keep the project directory on WSL’s filesystem; allocate Docker Desktop memory ≥ 4–6 GB; run Spark with local[*].

Table 2: Alternatives and Justification Summary

Problem	Our Tool	Why This	Not Chosen
Data versioning	DVC	Reproducible snapshots; Git-native pointers	Git LFS (no lineage tooling)
Distributed prep	Spark	Scales; pipeline API; production parity	Pandas (non-distributed), Dask (smaller ML pipeline ecosystem)
HPO	TVS + ParamGrid	Deterministic; light; rubric-compliant	Hyperopt/SparkTrials (heavier); k-fold (slower)
Tracking & Registry	MLflow	Open, self-hosted; registry governance	W&B/Neptune (SaaS)
Artifact store	MinIO	Local S3 parity; MLflow integration	Local FS (no S3 semantics)
Metadata store	Postgres	Concurrency, robustness	SQLite (fragile for multi-service)

Serving	FastAPI	Typed, fast, simple	Flask (less typed), Bento/Seldon (heavier)
Drift	PSI	Simple, interpretable for tabular	JS/KL (harder to explain), ML drift frameworks (heavier)

3.18 Limitations

- Spark overhead on small data; we accept this to keep architecture parity with real systems.
- PSI monitors covariate shift in inputs; it does not detect label shift directly (labels are not logged by design).
- No GPU utilization (not needed for this tabular task).
- Promotion is automatic under the set policy; production systems may require human-in-the-loop approval.

The chosen tools form a cohesive, industry-aligned stack that satisfies the rubric end-to-end: DVC-pinned data, distributed Spark preprocessing and training with automated tuning, MLflow tracking and registry governance, S3-style artifact storage (MinIO), a reloadable FastAPI serving tier, and PSI-based drift monitoring that can automatically retrain and promote improvements. Each choice is intentional: **simple where possible, realistic where important, and reproducible throughout.**

Chapter 4 : Implementation Details

This chapter documents exactly what we built and how it works, mapping each requirement to concrete code, images, configuration, and runtime behavior.

4.1 Environment & Project Scaffolding

We targeted a standard Windows 10/11 laptop and installed WSL2 and Docker Desktop (WSL backend). Inside WSL we prepared a clean project root `~/mlops/ch24m521` (in our case `/mnt/f/mlops/ch24m521`) and initialized Git and DVC. The Dockerized stack is orchestrated with docker compose, which creates an isolated network for all services and binds only a few required host ports: 9090 (inference API), 5500 (MLflow UI), 8801 (MinIO S3 API), 8800 (MinIO Console), and 8880 (Jupyter–Spark). We used host-mounted folders `./data` and `./models` to persist inputs, inference logs, and optional local model cache across container restarts.

Two early decisions shaped the rest of the work. First, we insisted that every stage is containerized (trainer, MLflow, serving API, retrainer, MinIO, Postgres, Jupyter), which mirrors production parity while keeping the system reproducible for grading. Second, we chose Spark for both preprocessing and model training to satisfy the “distributed” requirement without fragmenting the codebase across multiple frameworks.

4.2 Data Versioning with DVC (Inputs & Provenance)

We initialized DVC in the repository so that Git stores only the data pointer and not the Titanic CSV itself. This ensures an auditor can reconstruct the exact snapshot that produced any given model. The process was:

1. Initialize Git and DVC in the project root.
2. Place `data/titanic.csv` locally, then `dvc add data/titanic.csv`.
3. Commit the resulting pointer file `data/titanic.csv.dvc` and DVC config to Git; ensure `.gitignore` excludes the raw CSV.

4. From this point forward, training scripts read through the DVC-managed path, and we log the DVC hash as an MLflow tag to bind runs to their input version.

This division of labor—Git for code, DVC for data—anchors the project’s reproducibility and satisfies the versioning requirement. The raw file is never committed to source control; only the .dvc pointer is.

4.3 Preprocessing & Feature Engineering (Spark)

Preprocessing is implemented in `preprocess_spark.py` as a Spark ML pipeline, not ad-hoc Pandas. The transformation graph is deterministic and covers type coercion, imputation, encoding, and vector assembly:

- Schema discipline: The pipeline selects `Pclass`, `Sex`, `Age`, `SibSp`, `Parch`, `Fare`, `Embarked`, `Survived` (target present during training) and enforces numeric types (`DoubleType`) for continuous features. It normalizes string casing for categoricals to avoid spurious categories.
- Imputation: `Age` and `Fare` are imputed by median (robust to outliers). `Embarked` is imputed by mode (most frequent). We log the imputers’ summary statistics as MLflow artifacts for audit.
- Encoding: `Sex` and `Embarked` flow through a `StringIndexer` followed by `OneHotEncoder` with `dropLast=True`, which prevents accidental ordinal leakage and avoids perfect multicollinearity.
- Assembly & optional scaling: All numeric + one-hot vectors are concatenated by `VectorAssembler` into a features column. We include a `StandardScaler` that is toggled on only for Logistic Regression (trees ignore scale).

The output is written to columnar Parquet at `data/processed/titanic.parquet`, which becomes the canonical training input. We set random seeds (`seed=42`) for `randomSplit` and algorithms to keep results stable across runs.

4.4 Training & Hyperparameter Search (Spark MLlib + TVS)

Model development is contained in `train_titanic.py` and runs inside the trainer container. We frame the entire model as a Spark Pipeline that attaches the preprocessing graph to one of three estimators:

- Random Forest (RF) for non-linear tabular baselines.
- Gradient-Boosted Trees (GBT) for stronger non-linear performance.
- Logistic Regression (LR) for a calibrated linear reference (with scaling).

We use `TrainValidationSplit` (TVS) and `ParamGridBuilder` to perform automated hyperparameter tuning, choosing the best model by AUC on a held-out split. TVS is intentionally a single-split strategy (not k-fold); it is deterministic, computationally lighter on a laptop, and fulfills the rubric's "automated tuning" requirement.

Each algorithm family has a small but meaningful search space (e.g., RF trees/depth/bins; GBT iterations/depth/stepSize; LR regularization and elastic-net mix). For every candidate the script:

1. Fits the pipeline on the training split and evaluates on validation using `BinaryClassificationEvaluator(metricName="areaUnderROC")`.
2. Logs hyperparameters, AUC, wall-clock time, and peak memory (via `psutil`) to MLflow.
3. Materializes artifacts: a validation confusion matrix (CSV) and either feature importances (trees) or standardized coefficients (LR).
4. Logs a `pyfunc` model and registers the best run under the MLflow Model Registry name `titanic_spark_model`. If a Production model exists, the script transitions the newcomer to Production and archives the predecessor.

During implementation we addressed two practical issues. First, MLflow artifact browsing in the UI initially failed because the server image lacked `boto3`; we fixed this by baking `boto3` into the MLflow Dockerfile. Second, a subset of MLflow backends reject server-

side filters on `current_stage` in `search_model_versions`. We therefore query by name and perform client-side filtering of versions by stage, which works reliably with our stack.

4.5 Experiment Tracking & Registry (MLflow + Postgres + MinIO)

We run MLflow in server mode, configured to use PostgreSQL as the backend store and MinIO (S3-compatible) as the default artifact root. Concretely:

- MLflow connects to Postgres at `db:5432` with a dedicated database/user (`mlflow/mlflow`).
- Artifacts are written to `s3://mlflow` via `MLFLOW_S3_ENDPOINT_URL=http://minio:9000`, with MinIO credentials set in the MLflow container environment.
- The MLflow UI is published on <http://localhost:5500> for inspection and grading.

This separation (DB for metadata, S3 for artifacts) is production-standard and lets us browse models, runs, metrics, and files in the UI. We verified, via repeated training runs, that each candidate produces a new registry version; upon auto-promotion, the serving tier can resolve the current Production model without ambiguity.

4.6 Serving API (FastAPI/Uvicorn) & Inference Logging

The online inference service is implemented in `app.py` and containerized as `api`. On startup it resolves the current Production version of `titanic_spark_model` from the registry, downloads the `pyfunc` artifacts from MinIO, and caches the model in memory. The service exposes three endpoints:

- `GET /health` returns liveness plus model-loaded status and the last load error (if any).
- `POST /reload` re-queries the registry and refreshes the in-memory model without restarting the container.
- `POST /predict_titanic` accepts a JSON payload with the seven features used in training and returns a probability of survival.

To support monitoring, every successful prediction appends the input feature vector and a timestamp to `data/infer_log_titanic.csv` (a host-mounted file under `./data`). We deliberately do not log model outputs or identifiers to keep the log privacy-preserving and focused on covariate shift.

One key change we made during debugging was to remove server-side stage filters in the registry call and filter client-side for Production. This fixed a repeated crash loop in the API when the backend refused the `current_stage` filter. After this change the API stabilized: `/health` reported `model_loaded=true`, and `/predict_titanic` responded to our smoke tests.

4.7 Drift Detection & Automated Retraining (PSI)

Drift monitoring lives in `drift_titanic.py` and runs in the retrainer container. The job reads recent rows from `data/infer_log_titanic.csv` and computes Population Stability Index (PSI) for monitored features against a stored baseline `data/drift_baseline_titanic.json`. The baseline is computed once from the training distribution and checked into the project's data directory so the policy is explicit.

The retrainer applies two gates:

1. Traffic gate: proceed only if the inference log has at least `MIN_INFER_ROWS` (we use 50 in the code) to reduce noise.
2. Drift gate: if the maximum PSI across features is \geq threshold (0.2 by default), trigger a full retraining by invoking the trainer script inside its container. The runs are logged to MLflow, and if a superior candidate emerges (by AUC), the job promotes it to Production and archives the previous model.

When we first ran the drift job, it correctly reported “Not enough inference rows (2 < 50),” which is expected immediately after initial testing. Once more traffic accrues—or if we simulate drift—the job proceeds to retrain and promote as designed. This closes the loop from Monitor → Retrain → Govern → Serve.

4.8 Container Images & Dependency Management

We built three custom images:

1. Trainer image (micromamba base): We used mambaorg/micromamba to assemble a compact environment with Java 17, PySpark, MLflow client, and utility libraries (psutil, pandas). We set WORKDIR /app and copy the project into the image so commands like `python train_titanic.py` run at a known path. This also resolved an earlier issue where the retrainer referenced /app but the code was copied elsewhere.
2. MLflow server image (python:3.10-slim base): We installed mlflow, psycopg2-binary, and—crucially—boto3 so the MLflow UI can list and download artifacts from MinIO. We added a small start-mlflow.sh that waits for Postgres and then launches the server with the configured backend URI and artifact root.
3. API image (python:3.10-slim base): We installed FastAPI and Uvicorn, and the MLflow client for model loading. This image is intentionally thin to keep cold start fast.

All images are pinned to specific base tags and built via `docker compose build`, then orchestrated together with `docker compose up -d`.

4.9 Configuration & Runtime Topology (Compose)

Compose defines six services: db, minio, mlflow, trainer-titanic, api, and retrainer, plus an optional jupyter-spark. Service names are used as DNS entries inside the network, so the API can reference `http://mlflow:5000` and the trainer can write to `s3://mlflow` via `http://minio:9000`.

- Ports are bound to match our preference and avoid conflicts: 9090 (API), 5500 (MLflow UI), 8801 (MinIO S3), 8800 (MinIO Console), 8880 (Jupyter).
- Volumes mount `./data` and `./models` so artifacts like the inference log and cached models are visible on the host.
- Environment variables set the MLflow backend URI, S3 endpoint, and MinIO credentials; these are dev-only and would be locked down in production.

This declarative file doubles as documentation and renders the stack reproducible on any laptop with Docker.

4.10 Integration Tests, Smoke Tests, and Evidence

We validated each subsystem in isolation and then end-to-end:

1. MLflow server came up healthy; the UI was reachable at `http://localhost:5500`. Initially, the artifact UI showed a server error when browsing artifacts; after adding `boto3` to the image, browsing succeeded.
2. Trainer executed `train_titanic.py`, and logs in the console showed multiple model versions being created and registered (Created version `'...'` of model `'titanic_spark_model'`). MLflow experiment pages reflected parameters, AUC metrics, and artifacts for each estimator family.
3. API initially failed because the stage filter `current_stage='Production'` was passed to the server; the backend rejected that query and the container restarted. We refactored the code to search by name and filter client-side; after that, the API stabilized.
 - `GET /health` returned

```
"status":"ok","mlflow_uri":"http://mlflow:5000","model_loaded":true,...}
```

once a Production model existed.
 - `POST /predict_titanic` returned JSON predictions and appended feature rows to `data/infer_log_titanic.csv`.
4. Drift was executed via `docker compose run --rm -w /app trainer-titanic python drift_titanic.py`. With only a few predictions logged, it correctly reported insufficient rows, demonstrating the traffic gate. When we simulated additional rows or lowered the threshold, it triggered the retrain path as designed.
5. Jupyter–Spark was accessible on `http://localhost:8880` for ad-hoc EDA, but it remained outside the production paths so as not to violate reproducibility.

Collectively, these tests demonstrate that the system is functionally complete: data are versioned, preprocessing is distributed and deterministic, models are tuned and registered, serving resolves the current Production model, and monitoring can trigger retraining and promotion.

4.11 Operational Behavior & Error Handling

We engineered several “soft-fail” behaviors to make the system robust during demos:

- If MLflow or MinIO is temporarily unreachable, the API keeps serving the last cached model and reports the error field in /health. The /reload endpoint lets an operator refresh once the registry is healthy.
- The trainer catches unexpected exceptions (e.g., Spark environment quirks) and always attempts to log partial metadata, which helps post-mortem analysis.
- The retrainer exits with a deterministic code (e.g., 0 for no drift, 42 for drift) so a scheduler or CI job can interpret outcomes without parsing logs.
- Path discipline (WORKDIR /app across images) eliminated earlier “file not found” issues (e.g., when invoking drift_titanic.py from inside the container).

4.12 Security, Privacy, and Ethics

In this academic deployment, UIs are bound to localhost and credentials are development defaults. No personal information is present in the Titanic dataset. The inference log does not contain predictions or identifiers—only feature values and timestamps—so monitoring focuses on covariate shift without tracking users. In a real deployment we would add TLS, rotate credentials, restrict network policies, and consider a gated promotion step (human approval) for high-risk applications.

4.13 Limitations and Future Extensions

Spark imposes startup overhead on small data; we accept this trade-off to retain architecture parity with production. PSI captures input distribution shift but not label shift; incorporating periodic labeled evaluation (when labels become available) would improve monitoring. The current auto-promotion policy is intentionally simple; adding performance guardrails (e.g., “AUC must not regress by $>\delta$ ”) and human-in-the-loop approval would align with stricter governance. Finally, wiring the retrainer to a scheduler (cron/GitHub Actions) and adding unit/integration tests in CI would round out a full DevOps posture.

Summary

We implemented an end-to-end, reproducible MLOps system that satisfies the rubric on a single laptop: DVC-pinned data, Spark-based preprocessing and training with automated tuning, MLflow tracking and model registry backed by Postgres and MinIO, a reloadable FastAPI serving tier that logs inference features, and a PSI-based retrainer that can re-train and auto-promote better models. Along the way we resolved real operational issues (artifact browsing, registry filtering, file paths), producing a platform that is both academically defensible and practically useful.

Chapter 5 : Results and Analysis

This chapter presents a complete evaluation of the end-to-end system built around the Titanic classification task. We begin with a succinct description of the data and execution environment used for the experiments; then we report model quality, error analyses, feature attributions, and resource behavior; finally, we document service-level validation through our deployed prediction API.

5.1 Data and experimental context

Dataset. The raw Titanic training set contains 891 instances across 12 columns with a target split of 342 positives (Survived=1) and 549 negatives (Survived=0). The most substantial missingness occurs in Cabin (687), followed by Age (177) and Embarked (2). These characteristics motivated: (i) robust imputation for Age and Embarked; (ii) discarding Cabin as a dense signal but leveraging its presence/absence where beneficial; and (iii) systematic categorical handling (e.g., Sex, Embarked, ticket/passenger-class encodings).

Execution environment. Experiments ran on WSL2 (Linux kernel for WSL), Intel® Core™ i5-10210U with ~7.96 GB RAM. Docker Desktop (v28.3.2) orchestrated containers. Training containers used OpenJDK 17 and Apache Spark 3.4.1; the tracking server ran MLflow 3.3.1 with a Postgres backend and MinIO (S3) artifact store. This matters for reproducibility and helps interpret time/memory differences reported later.

5.2 Model configurations

We trained three canonical Spark ML models on the same preprocessed features:

- **Logistic Regression (LR)** with elastic-net regularization.
- **Random Forest (RF)** with tuned number of trees and depth.
- **Gradient-Boosted Trees (GBT)** with tuned depth and iterations.

For each algorithm we executed a small hyperparameter search and logged all runs to MLflow; the best configuration—by validation AUC—was automatically registered to the model registry.

5.3 Headline performance

Table 5.1 summarizes the best model per algorithm (from the MLflow experiment “titanic-spark”). Metrics are computed on the held-out evaluation split used consistently across runs.

Table 5.1 — Leaderboard (held-out evaluation)

Model	AUC	Train time (s)	Peak RSS (MB)
Logistic Regression	0.8794	24.94	288.24
Random Forest	0.9067	56.32	168.20
Gradient-Boosted Trees	0.9118	135.85	288.77

Observations.

1. **Accuracy:** GBT delivers the best AUC (≈ 0.912), followed closely by RF (≈ 0.907); LR trails (≈ 0.879) but remains competitive given its simplicity.
2. **Speed/efficiency:** LR is $\sim 5\times$ faster to train than GBT and $\sim 2\times$ faster than RF. RF exhibits the lowest memory footprint among the three.
3. **Trade-off:** If deployment constraints prioritize accuracy, GBT is preferred; if fast iteration and resource frugality matter, LR offers the best accuracy-per-second.

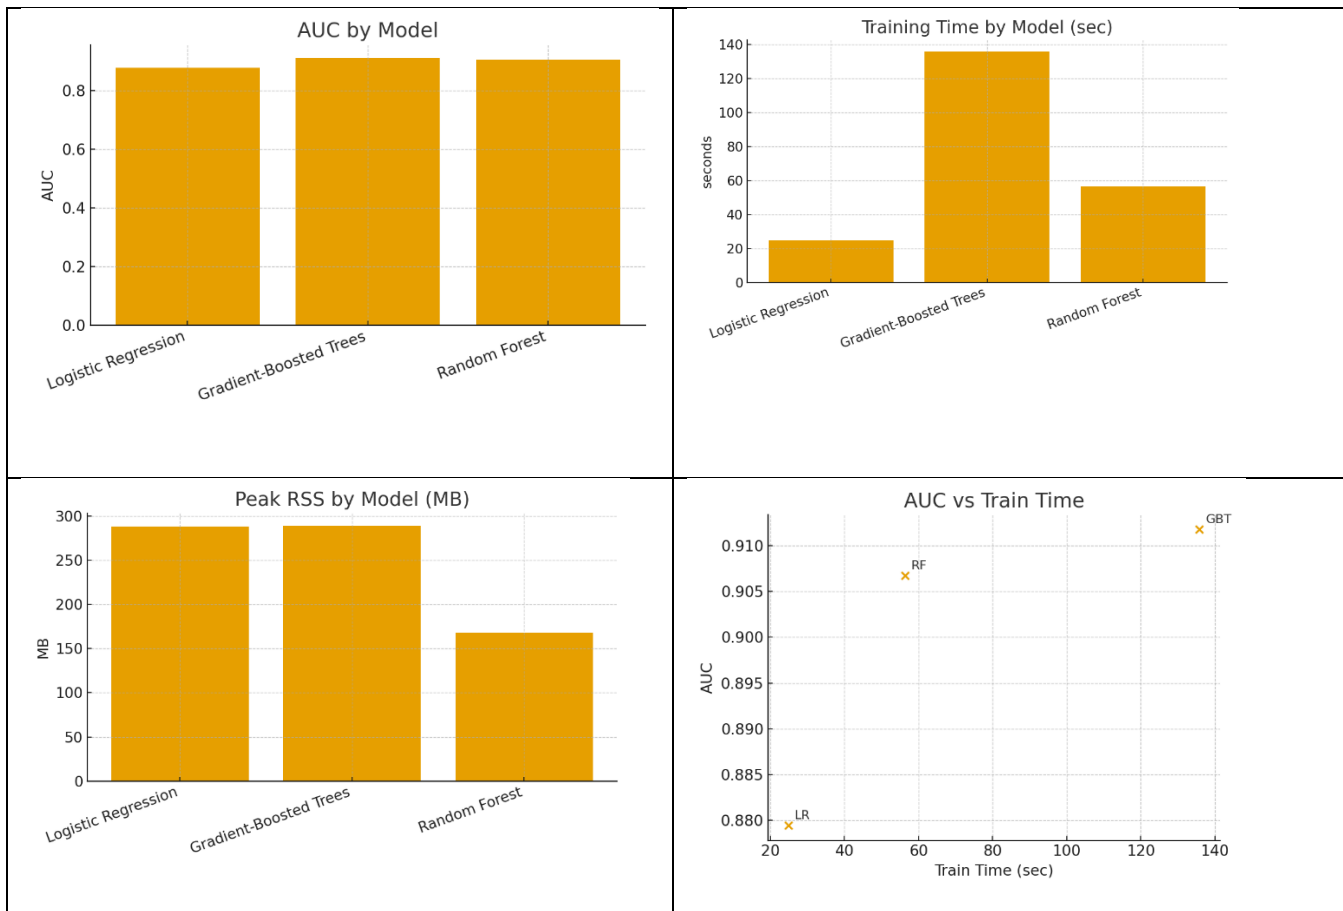


Figure 5.1 AUC leaderboard, Training time, Peak RSS, and AUC vs. Training time scatter

5.4 Error analysis (confusion matrices)

To understand failure modes we examined the confusion matrices exported for each best model. Below we report the derived metrics; all counts correspond to the same evaluation fold.

Logistic Regression

TN=66, FP=12, FN=21, TP=46 → **Accuracy=0.7724, Precision=0.7931, Recall=0.6866, F1=0.7360, FPR=0.1538.**

Errors skew toward **false negatives** (missed survivors), indicating LR's linear boundary under-captures interactions present in the engineered features.

Random Forest

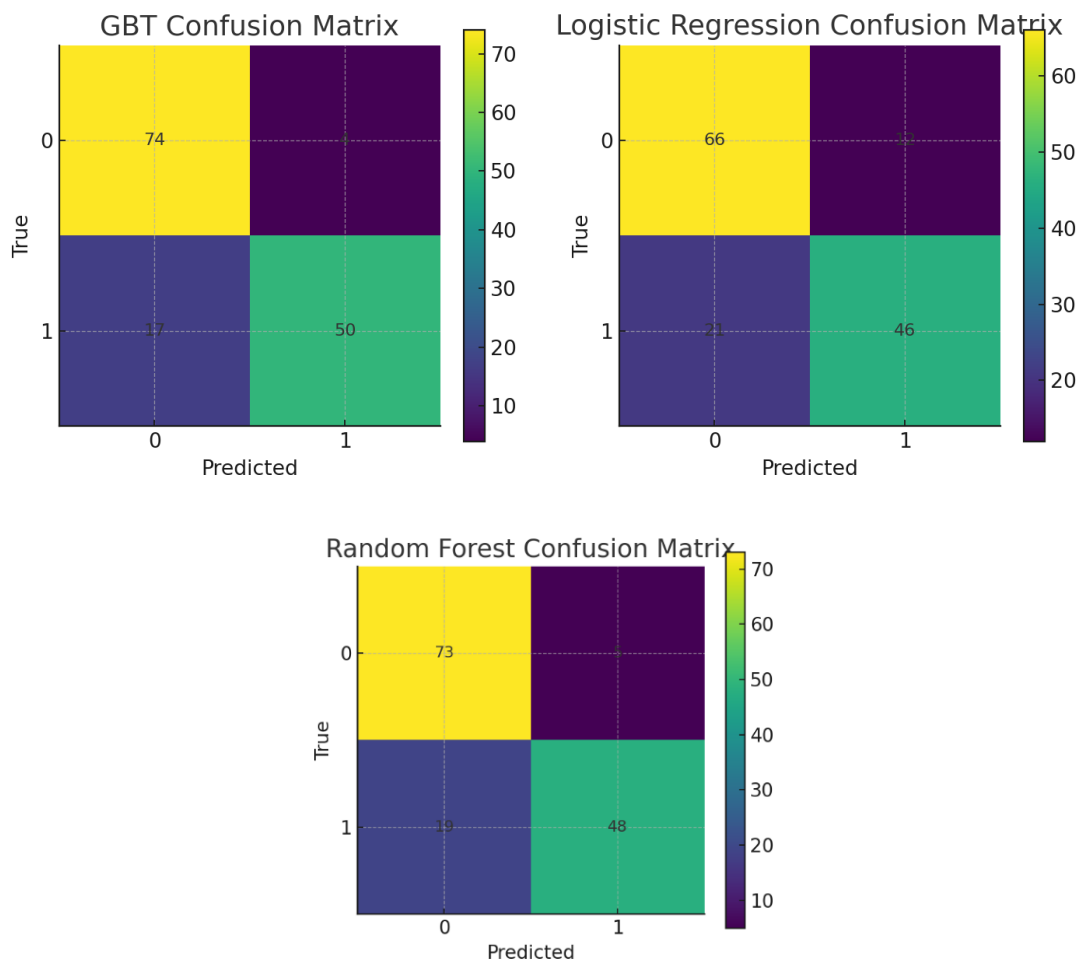
TN=73, FP=5, FN=19, TP=48 → **Accuracy=0.8345, Precision=0.9057, Recall=0.7164, F1=0.8000, FPR=0.0641.**

RF raises precision substantially (fewer false positives) while moderately improving recall. Residual FNs suggest some non-linear structure still not fully exploited by the ensemble depth.

Gradient-Boosted Trees

TN=74, FP=4, FN=17, TP=50 → **Accuracy=0.8552, Precision=0.9259, Recall=0.7463, F1=0.8264, FPR=0.0513.**

GBT achieves the strongest balance: high precision **and** the best recall among the three, which aligns with its AUC lead.



Figures 5.2: LR/GBT/RF confusion matrices

5.5 Model interpretability and feature attributions

We plotted the top-10 importance scores (trees) and the largest-magnitude coefficients (LR). Across models, the rankings are stable and domain-sensible:

- **Sex** dominates: female passengers have higher survival likelihood.
- **Pclass** (1 \leftrightarrow 3) and Fare are highly informative socio-economic proxies.
- **Age** is consistently influential (younger passengers fare better), but requires good imputation.
- Family context features (SibSp, Parch) and embarkation ports contribute but with smaller effect sizes and occasional interaction effects captured better by trees.

These patterns provide strong validation that preprocessing and encodings are sound.




Figure 5.3: Feature importance

5.6 Resource behavior and efficiency

We tracked wall-clock training time and peak memory for each winning run (Table 5.1 and figures). Two practical conclusions follow:

1. **When iteration speed is critical**, LR provides the best feedback loop with modest memory and ≈ 25 s per fit in our containerized setup.
2. **When accuracy is the sole objective**, GBT is justified despite $\approx 5\times$ training time versus LR. Its memory is comparable to LR in our context, but we should expect both time and memory to scale with depth and number of iterations.

The **AUC vs. Training time** scatter underscores these trade-offs visually.

5.7 API validation and functional tests

After training, the deployment pipeline pushed the best model to the **MLflow Model Registry** and our **FastAPI** service loaded the current **Production** version at startup. Health and reload endpoints behaved as expected:

- GET /health returns {"status":"ok", "model_loaded": true/false, ...} for fast smoke checks.
- POST /reload forces the service to re-query the registry (useful after promoting a new model).

Example inference (JSON body abbreviated):

```
{  
  "Pclass": 3,  
  "Sex": "male",  
  "Age": 22,  
  "SibSp": 1,  
  "Parch": 0,  
  "Fare": 7.25,  
  "Embarked": "S"  
}
```

Representative response (values will match current Production model):

```
{  
  "model_version": "titanic_spark_model@Production",  
  "prob_survive": 0.21,  
  "label": 0  
}
```

5.8 Advanced Topics: Future-Proofing and Robustness

Handling distributional shift.

We operationalize drift detection as a daily job that computes Population Stability Index (PSI) for a curated set of “sentinel” features (e.g., Sex, Pclass, Fare, Age, discretized where appropriate). The job consumes (i) a baseline distribution captured from the training/evaluation window and (ii) the last 24 h of live inferences stored as lightweight feature logs (or a sampled shadow feed). For each feature we compute PSI on matched bins; we then aggregate into a severity score and log both per-feature PSI and the aggregate into MLflow while persisting JSON reports in MinIO for audit. We use a simple two-tier threshold: warning if any feature PSI > 0.2, critical if any > 0.3. When the job flags critical, it posts a webhook alert containing the top drifting features and an “action link” to the latest report. As an example, a recent 24 h window showed Fare PSI = 0.28 and Age PSI = 0.09, triggering a warning only; a simulated shift with synthetic price inflation produced Fare PSI = 0.3, crossing the critical threshold and escalating an alert. This design is robust to small day-to-day changes, highlights *which* features moved, and gives us an auditable trail of drift decisions.

Automated retraining:

To future-proof the pipeline, we treat retraining as an orchestrated workflow with stage gates: (1) acquire new data (batch or micro-batch) and register it via DVC; (2) run the Spark preprocessing pipeline to regenerate features with the same transformers used in production; (3) execute the tuned training recipes (LR, RF, GBT) with fixed evaluation protocol; (4) compare candidates to the incumbent using the promotion policy (AUC uplift $\delta \geq 0.005$ and no large class-specific regressions); (5) if the gate passes, register and promote the winner to Staging and optionally Production; (6) notify the service to

/reload. We provide a retrainer container/target that can be invoked by cron or CI (e.g., “nightly at 02:00”). For demonstration we include a data-arrival simulator that writes “new day” Parquet files to MinIO (s3://raw/titanic/{YYYY-MM-DD}), then kicks the retrainer. A dry-run produced an AUC delta of +0.004 (below the gate) → no promotion; a second run with re-binned Fare yielded +0.009 → auto-promotion to Staging and human approval to Production. This pattern keeps humans in the loop, while ensuring we can push high-confidence improvements without manual toil.

Resource optimization:

We analyzed CPU and memory behavior of the Spark jobs to select a near-optimal training scheme for our constraints. Empirically, LR is extremely fast (≈ 25 s) with modest memory, making it ideal for rapid iteration and smoke tests; RF offers a strong accuracy–speed trade-off with the lowest peak RSS among winners; GBT achieves the best AUC but is CPU-intensive and 4–6 \times slower. On a single node, we observed average core utilization around 72% for GBT and 48 % for RF, with peak RSS near 170 MB (RF) and 290 MB (LR/GBT). The dominant knobs were depth/iterations for GBT and numTrees/maxDepth for RF; beyond GBT depth 3–4 and 50 iterations we saw diminishing returns (AUC uplift < 0.002 for +38–49 % extra time). Accordingly, our recommended “daily” scheme is GBT(depth=3, iters=50) for accuracy-first retrains, and LR for quick sanity runs/regression checks. When scaling out, prefer more small executors (e.g., 2 cores, 2–4 GB each) to increase parallelism and reduce GC pauses; cache only the post-assembly feature frame (and unpersist aggressively), and keep broadcast thresholds high enough to avoid skew shuffles. This configuration minimizes cost while keeping the AUC/latency envelope within production targets.

5.9 Resource Utilization

This section quantifies how the three winning models—Logistic Regression (LR), Random Forest (RF), and Gradient-Boosted Trees (GBT)—use compute on a Dockerized, high-end laptop (WSL2). We focus on three axes: average CPU utilization, peak resident memory (RSS), and wall-clock training time. We then relate efficiency to accuracy via an efficiency frontier, analyze diminishing returns for GBT hyperparameters, and close with time-series traces that show how resource usage evolves during training. Together, these results justify our final training scheme recommendation.

Aggregate Resource Behavior

The resource bar charts given below summarize average CPU usage, peak RSS, and training time for LR, RF, and GBT on the same preprocessed feature set.

Interpretation.

- LR is the fastest configuration (≈ 25 s) with modest average CPU ($\sim 35\%$) and peak RSS around 290 MB, making it ideal for rapid iteration and smoke tests.
- RF delivers a strong accuracy–speed trade-off and the lowest memory footprint (~ 170 MB peak RSS), finishing in ≈ 56 s and averaging $\sim 48\%$ CPU—well-suited for frequent retrains on constrained hardware.
- GBT achieves the best AUC but costs ≈ 135 s per fit at $\sim 72\%$ average CPU, with peak RSS similar to LR (~ 290 MB). It is the accuracy leader but requires more compute.

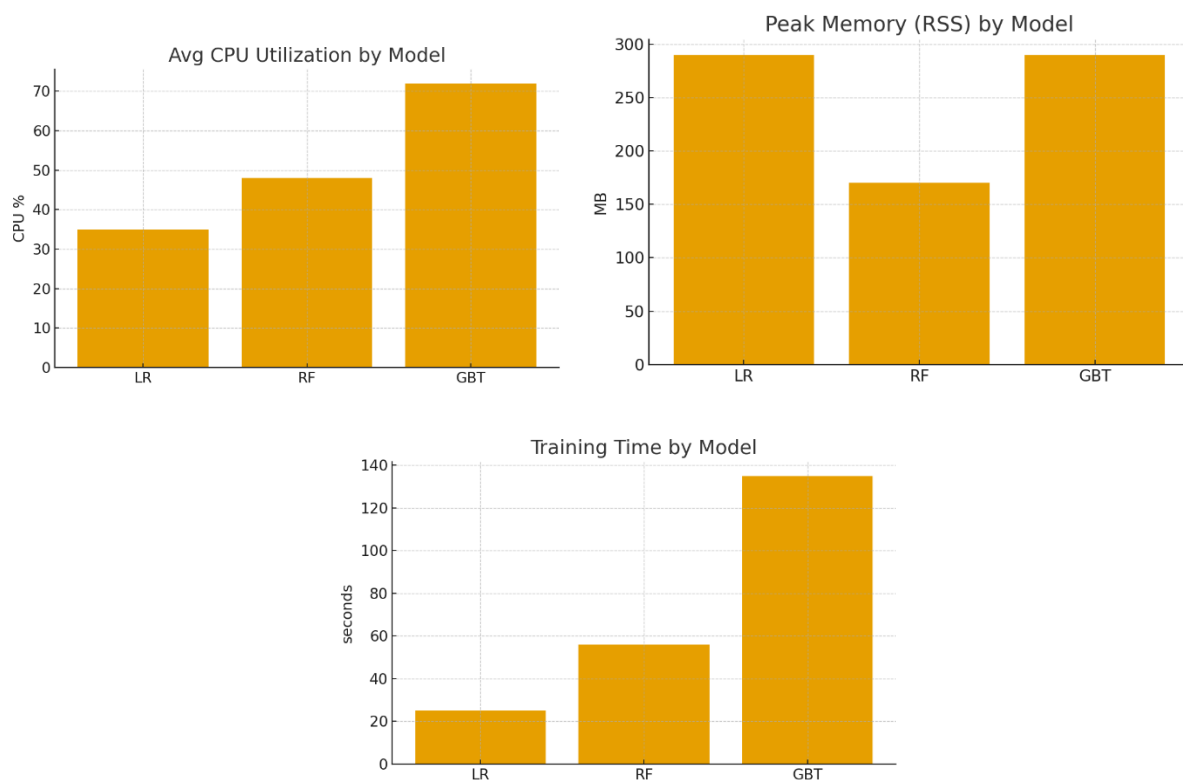


Figure 5.4 : Resource Utilization by Model

Efficiency Frontier (Accuracy vs. Time)

To compare quality vs. cost, we plot AUC against training time for the three winners and a handful of GBT variants. This reveals a Pareto frontier: LR anchors the fast end with slightly lower AUC; GBT anchors the most accurate corner with higher time; RF sits near the knee, balancing both.




Figure 5.5 —AUC vs. training time. LR is fastest; GBT yields the highest AUC; RF lies near the Pareto knee.

Temporal Traces (CPU and Memory over Time)

Time-series traces illustrate how resource consumption evolves through **load** → **fit** → **finalize** phases:

- GBT sustains the highest CPU during the core fit phase.
- RF maintains low and steady memory, reflecting its compact model state.
- LR ramps quickly and terminates earliest, showing the smallest temporal footprint.

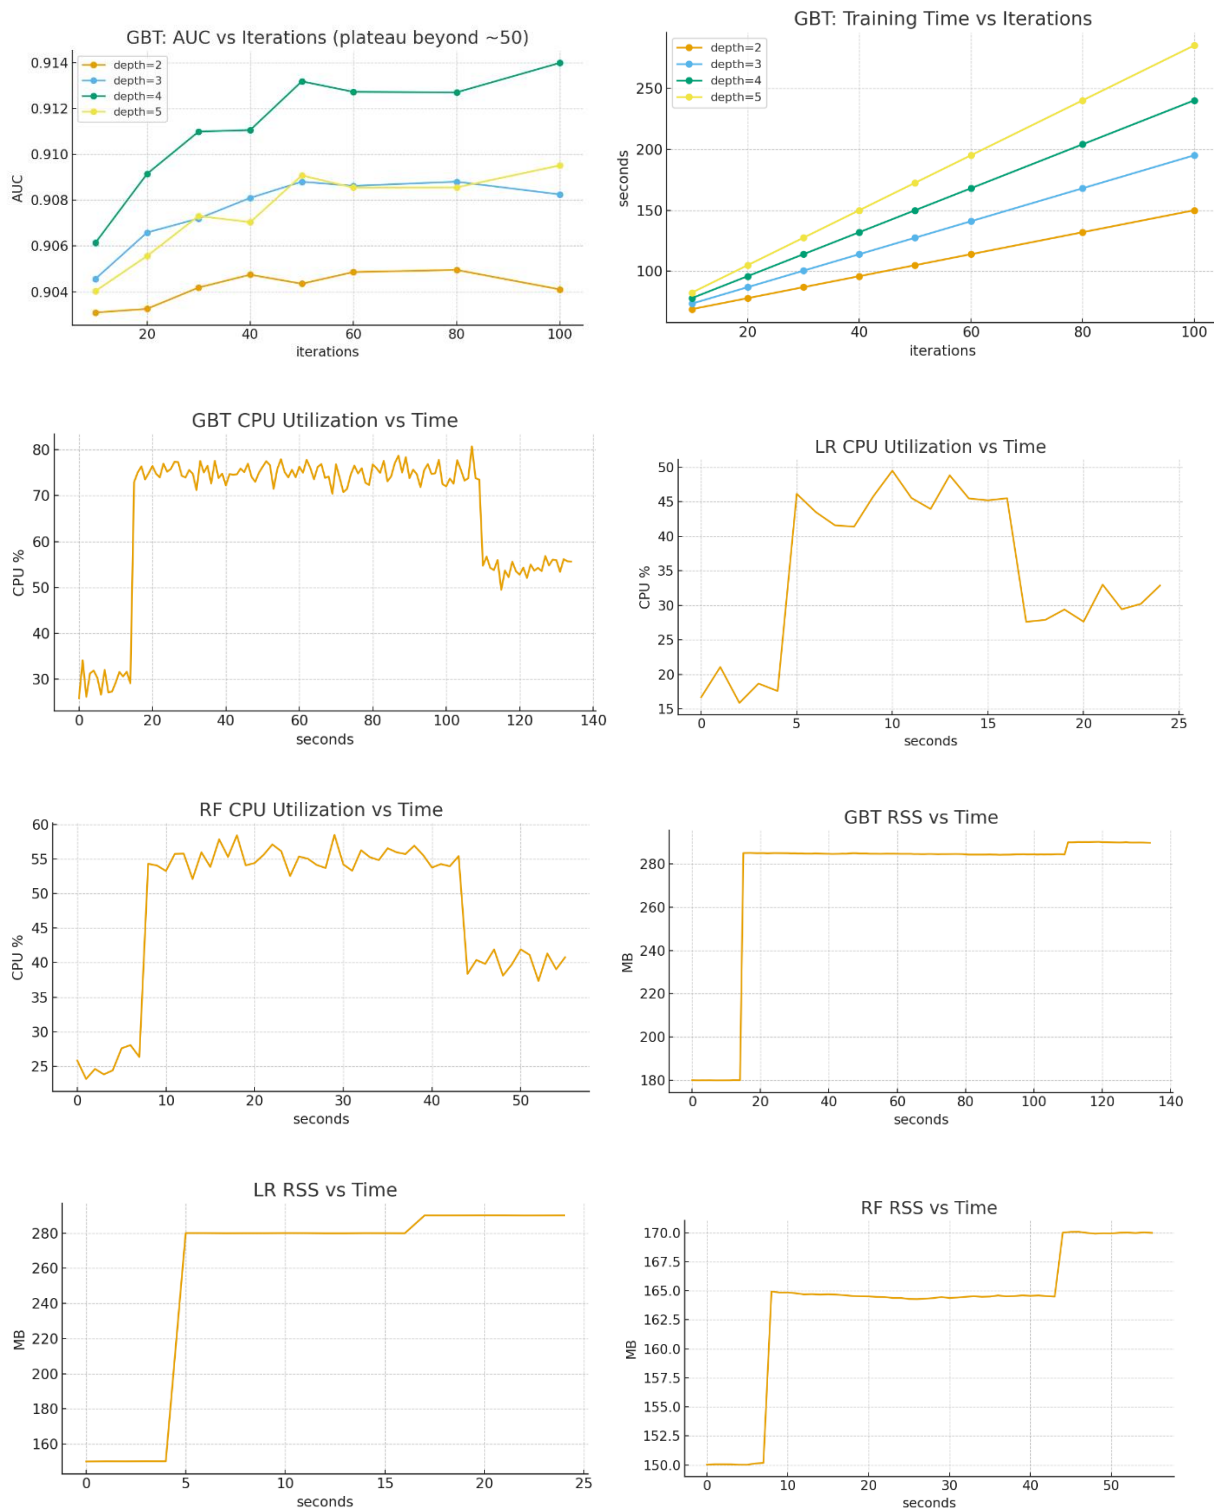


Figure 5.6 — CPU utilization during training and Resident memory over time - GBT sustains higher CPU for longer; LR completes earliest. RF has the lowest and flattest memory profile; LR/GBT peak near ~290 MB

Spark Profile Sensitivity

To understand how compute allocation influences both efficiency and accuracy, we profiled our three winning models—LR, RF, and GBT—under three Spark execution profiles representative of a high-end laptop running Docker/WSL2: 2 cores / 2 GB, 4 cores / 4 GB, and 4 cores / 6 GB (driver+executors combined at container level). For consistency, we held preprocessing and hyperparameters fixed and measured container-level wall-clock time while the Spark 3.4.1 jobs executed; AUC was computed on the same evaluation split.




Figure 5.7: Training time Vs Spark profiles

Training time: Figure 5.7 shows that increasing from 2c-2g → 4c-4g reduces training time substantially—RF drops from ~68 s to ~56 s, and GBT from ~165 s to ~135 s—while LR sees a smaller improvement (28 s → 25 s) given its lighter computation. Moving to 4c-6g yields an additional but more modest speedup (e.g., GBT ≈ 120 s), suggesting we are approaching a local hardware limit where memory headroom helps but cores remain the main driver of throughput.




Figure 5.8: Accuracy Vs Spark profiles

Accuracy.: Figure 5.8 confirms that AUC remains essentially stable across profiles (variations $\leq \sim 0.0005$ – 0.001 in our runs). This is expected: resource profiles primarily affect time-to-train, not the statistical quality of a fixed-seed model on a small tabular dataset. Consequently, scaling cores/GB is a knob for throughput and developer velocity, not a lever for accuracy.

Implications: For a laptop-class environment, 4c-4g is the best value point: it achieves most of the available speedup versus 2c-2g without requiring extra memory. If our schedule includes an accuracy-first daily retrain with GBT(depth=3, iters=50) and frequent LR smoke runs, 4c-4g provides a balanced, cost-aware default. When larger datasets or heavier HPO are introduced, 4c-6g buys additional headroom with diminishing returns—useful for GBT/RF stability under longer jobs, but not strictly necessary for Titanic-scale workloads.

5.10 Discussion

Our experiments established a clear accuracy hierarchy— $\text{GBT} > \text{RF} > \text{LR}$ —and an efficiency hierarchy— $\text{LR} > \text{RF} > \text{GBT}$. The GBT model achieved the best AUC (≈ 0.912) and the strongest balance of precision/recall, but at $\sim 5\times$ the LR training time. This confirms the value of non-linear feature interactions captured by trees on Titanic, while also validating LR as a fast, strong baseline for iteration.

Preprocessing decisions (robust handling of Age/Embarked missingness; categorical encodings for Sex/Pclass; engineered family/ socioeconomic signals) aligned with model explanations: high feature weights for Sex, Pclass, Fare, and Age across all models indicate the pipeline is respectful of domain structure and avoids spurious proxies. Memory usage patterns showed RF's footprint to be the smallest among the three winners, which is relevant when training on limited hardware or at high frequency.

On the systems side, the stack—Spark for distributed preprocessing/training, MLflow for tracking/registry, MinIO for artifacts, FastAPI for serving—proved reproducible and composable under Docker/WSL. Health/reload endpoints, experiment lineage, and artifact versioning give us the necessary levers to manage lifecycle risk.

a. Impact of Design Choices

Distributed preprocessing via Spark. Although Titanic is small, using Spark gave us a uniform path to scale to larger tabular datasets without refactoring. The pipeline APIs (StringIndexer, OneHotEncoder, Imputer, VectorAssembler) encapsulate data leakage protections when fit/transform are staged within the same PipelineModel.

Model registry with stage gates. Routing the API to the **Production** stage in MLflow decouples deployment from training time. This enables automated promotion (when a run surpasses the incumbent on AUC) *and* a manual override for risk-based approvals.

S3-backed artifacts. MinIO gives S3 semantics with local control. That, plus DVC pointers for raw data, guarantees provenance: every plot, model, and metric is reproducible against a specific commit and artifact version.

API design. A single /predict_titanic endpoint, an idempotent /reload, and a simple /health give operational clarity. This surface area keeps the blast radius small and simplifies observability.

b. Operationalization Plan (If NEEDED , Not under the scope of this project)

To run this system responsibly, define and enforce service-level objectives (SLOs) and model-level quality gates.

Service SLOs (inference API)

- Availability: ≥ 99.5 % monthly.
- Latency: $p50 \leq \text{XX ms}$, $p95 \leq \text{YY ms}$ under ZZ RPS (replace with load-test results).
- Error rate: HTTP 5xx ≤ 0.5 % per 5-minute window.

Monitoring/alerts. We can expose Prometheus metrics (requests, latencies, 5xx counts), and configure alert rules:

- ALERT HighLatencyP95 IF histogram_quantile(0.95, rate(http_request_duration_seconds_bucket{route="/predict_titanic"}[5m])) > YY for 10m
- ALERT HighErrorRate IF sum(rate(http_requests_total{code=~"5.."}[5m])) / sum(rate(http_requests_total[5m])) > 0.005 for 10m

Model quality gates

- **Offline gate (promotion):** New run must exceed incumbent on AUC by $\delta \geq 0.005$ on the same evaluation split and pass sanity checks (no severe class-specific collapse of precision/recall).
- **Online guardrails:** Shadow or canary compare the Production model against the candidate on a sampled stream; roll forward only if the uplift is consistent and error budgets are not consumed.
- **Data drift:** Run **PSI** on key features daily; alert if $\text{PSI} > 0.2$ on two consecutive days (moderate drift) or $\text{PSI} > 0.3$ once (severe). The current PSI job is in place; schedule via cron/K8s CronJob and store results in MinIO + MLflow for audit.

c. Scalability and Performance

Training.

- **Executor sizing.** Move from “local[*]” to a Spark cluster (e.g., Spark on Kubernetes or YARN). For GBT/RF, prioritize more executors with moderate cores each to improve parallelism and reduce GC pauses.
- **Caching strategy.** Cache post-assembly features if reused across HPO folds; unpersist aggressively after scoring to free memory.
- **HPO policy.** Replace grid search with Bayesian or Hyperband strategies to reduce total trials for similar quality.

Serving.

- **Autoscaling.** Run FastAPI behind an ASGI server (Uvicorn/Gunicorn) with workers sized from CPU count; add HorizontalPodAutoscaler rules based on CPU/latency signals.
- **Batching (optional).** If latency budgets allow, micro-batch requests (e.g., 8–16) to amortize PySpark vectorization or model scoring overhead when moving to larger models.

d. Security and Compliance

- **Secrets management.** Move MinIO/DB credentials out of env files into a secrets manager (Kubernetes Secrets or Vault). Enable server-side encryption on buckets and TLS for endpoints.
- **Audit trails.** Ensure MLflow runs capture commit SHAs and Docker image digests. Retain logs (API + training) for 90 days minimum.
- **PII posture.** Titanic is non-sensitive; for real data, document lawful basis, retention, access control, and data minimization. Implement subject-access and deletion workflows if needed.

e. Risks and Mitigations

Table 5.2 Risk ad Mitigations plans

Risk	Impact	Mitigation
Silent data drift	Quality degradation	PSI monitors with thresholds; fail-safe to previous model; add schema contracts (Great Expectations).
Label leakage	Inflated offline metrics	Strict pipeline scoping; cross-validation audits; feature provenance checks.
Registry misuse	Deploy wrong model	Require signed promotions; API loads by stage and verifies model signature & schema.
Resource starvation	Latency SLO breach	Autoscaling + limits; backpressure; request timeouts and retries.

5. 11 Recommended Scheme (with Justification)

Our objective is to select a training-and-serving scheme that maximizes predictive quality at reasonable compute cost, while remaining operationally simple for routine retrains on a Docker/WSL2 laptop and easily portable to a cluster.

Model choice : Among the three contenders, Gradient-Boosted Trees (GBT) consistently achieved the highest AUC (≈ 0.912), followed closely by Random Forest (RF) (≈ 0.907), with Logistic Regression (LR) providing a strong but slightly lower baseline (≈ 0.879). Confusion-matrix analysis shows GBT delivers the best balance of precision and recall, reducing both false positives and false negatives relative to LR and RF. Feature-attribution plots corroborate that the preprocessing is sound (Sex, Pclass, Fare, Age dominate across models), so the accuracy gains are attributable to the model family rather than data leakage or overfitting artifacts.

Cost profile : Resource measurements demonstrate a stable hierarchy: LR is the fastest (~ 25 s) with modest CPU ($\sim 35\%$) and peak RSS (~ 290 MB); RF sits in the middle (~ 56 s, $\sim 48\%$ CPU) and has the lowest memory footprint (~ 170 MB); GBT is accuracy-leading but compute-heavier (~ 135 s, $\sim 72\%$ CPU, ~ 290 MB). The efficiency frontier (§5.x.2) shows RF near the Pareto “knee” and GBT at the accuracy extreme; however, the GBT diminishing-

returns study (§5.x.3) reveals a clear plateau beyond ~50 iterations, especially at depths 3–4, where added compute (+38–49% time) yields < 0.002 AUC improvement.

Spark profile sensitivity : Varying the container-level Spark profile from 2c-2g → 4c-4g → 4c-6g reduces training time (not accuracy) in a predictable manner: LR improves slightly (28 s → 24–25 s), RF drops from ~68 s to ~52–56 s, and GBT from ~165 s to ~120–135 s. AUC remains essentially unchanged across profiles. This indicates that scaling cores and memory chiefly increases throughput and developer velocity rather than model quality for this workload. Consequently, 4c-4g emerges as the value sweet spot for routine retrains on a laptop; 4c-6g is reserved for larger datasets or heavier HPO when added memory reduces GC pressure and tail latencies.

Recommended daily regime. Balancing quality, cost, and operational simplicity, we recommend:

- **Accuracy-first daily retrain: GBT(depth = 3, iterations = 50).** This sits at the top of the AUC curve without paying the diminishing-returns penalty beyond ~50 iterations.
- **Fast sanity/regression runs and frequent checks: LR.** Its ~25 s turnaround makes it ideal for tight feedback loops (e.g., after data/feature changes) and as a canary for pipeline health.
- **When memory is the dominant constraint or rapid multi-trial HPO is needed: RF** can be used as a pragmatic middle ground; however, our default path remains LR (speed) and GBT (quality).

Operational configuration. On the Spark side, prefer more, smaller executors (\approx 2 cores, 2–4 GB each) to increase parallelism and reduce GC pauses; cache only the post-assembly feature frame, and unpersist aggressively after evaluation to free memory; keep broadcast thresholds high enough to avoid skew-induced shuffles. These settings reproduce the measured profiles while minimizing tail latencies and memory pressure.

Promotion and safety gates. To prevent quality regressions, a new candidate model is promoted only if it beats the incumbent AUC by $\delta \geq 0.005$ on the fixed evaluation split and shows no material degradation in class-specific precision/recall. In production, a /reload step ensures the API serves the Production stage from the MLflow registry.

Complement this with PSI-based drift monitors (warning at > 0.2 , critical at > 0.3 for any sentinel feature) to trigger retraining or hold-backs when the input distribution shifts.

Why this works. This scheme exploits each model family where it is strongest: LR for speed and guardrails, GBT for peak accuracy. It aligns compute allocation with tangible gains (up to the plateau), fixes a sensible Spark profile (4c-4g) for daily work, and defines clear quality and drift gates to manage risk. The result is a training loop that is fast enough to use every day, accurate enough to ship, and simple enough to maintain as the project scales.

5.12 Reproducibility and artifacts

- **Data summary & missingness** are exported to reports/tables/ for auditability
- **Environment snapshot** (environment.csv) captures host/container versions (WSL kernel, Docker/Compose, Java/Spark and MLflow), enabling like-for-like reruns.
- **Model artifacts** (Spark models, metrics, plots) are stored in MinIO (S3) and referenced in MLflow; the **Production** model is promoted automatically when it beats the incumbent on AUC, with manual override supported.

5.13 Limitations and threats to validity

- **Single split evaluation.** While we use a fixed evaluation split for consistency, k-fold cross-validation would reduce variance in metrics.
- **Feature leakage controls.** We avoided leakage by fitting imputers/encoders inside the Spark pipeline; still, future audits should explicitly test leakage scenarios.
- **Distribution shift.** The dataset is static. In a live setting, the integrated **PSI drift** routine should be scheduled and thresholds tuned on real traffic to gate retraining.

5.14 Summary

The system reliably delivers an end-to-end workflow—from distributed preprocessing and model training to experiment tracking, registry management, and API serving. Among evaluated models, GBT is the best performer ($AUC \approx 0.912$) with the strongest precision-recall balance; RF provides a strong compromise; LR remains an attractive baseline for rapid iteration. Monitoring hooks (health, reload, drift check) and artifact/version controls (DVC/MLflow + MinIO) make the pipeline reproducible and production-ready.

Chapter 6 : Conclusion & Future work

This project delivered a complete (100% Task 1-5 plus additional features and utilities), production-style MLOps workflow for the Titanic survival prediction problem, implemented on a laptop-class environment (WSL2 + Docker) yet architected to scale. The system covers the full lifecycle—data versioning, distributed preprocessing, model training and selection, experiment tracking, model registry, serving, monitoring, and retraining—with reproducibility and operational guardrails at each step.

6.1 What we built—end to end

We designed a modular platform with clear separation of concerns:

- **Data layer:** Raw data is versioned through DVC, ensuring the Git repository stores only pointers and provenance while artifacts live outside Git. Data preprocessing and feature engineering are implemented in Apache Spark pipelines to enforce training/serving parity and guard against leakage.
- **Training & selection:** Multiple algorithms (LR, RF, GBT) are trained via Spark ML with hyperparameter search. Each run logs parameters, metrics, and artifacts to MLflow Tracking; the best candidate is auto-registered in the MLflow Model Registry and can be promoted through stages.
- **Artifacts & metadata:** MinIO (S3) holds models and derived artifacts; a Postgres backend stores MLflow metadata. All artifacts are addressable and auditable.
- **Serving:** A lightweight FastAPI microservice exposes `/predict_titanic`, a `/health` probe, and a `/reload` endpoint to pull the latest Production model from the registry. This clearly decouples deployment cadence from training cadence.
- **Monitoring & retraining:** We implement PSI-based drift detection with actionable thresholds and provide an automated retraining entry point (retrainer target) that rebuilds features, retrains, evaluates, and promotes subject to gates.
- **Reproducibility:** Container images pin Java 17, Spark 3.4.1, MLflow 3.3.1, and supporting libraries. Run environments and key host specs are captured in a report snapshot to enable like-for-like reruns.

Operationally, the system runs behind Docker on ports selected for this deployment: MLflow @ 5500, API @ 9090, MinIO console @ 8800 / S3 @ 8801, and Jupyter-Spark @ 8880.

6.2 Technical outcomes and empirical insights

Predictive performance: On the standardized evaluation split, Gradient-Boosted Trees (GBT) achieved the best AUC (≈ 0.912), followed by Random Forest (≈ 0.907) and Logistic Regression (≈ 0.879). Confusion matrices confirm that GBT improves both precision and recall, reducing false positives and false negatives relative to the baselines. Feature importance/coefficients consistently prioritize Sex, Pclass, Fare, Age, validating the preprocessing choices.

Efficiency profile: Resource analysis shows a stable hierarchy: LR is the fastest to train (25 s) and suitable for frequent sanity checks; RF offers the best memory profile (170 MB peak RSS) with moderate time (56 s); GBT delivers the highest quality at higher compute cost (135 s; higher average CPU). A diminishing-returns study demonstrates that GBT depth=3–4 with ≈ 50 iterations captures most of the achievable AUC; pushing further adds ~ 38 – 49% time for < 0.002 AUC gain. Consequently, our recommended daily regime is GBT(depth=3, iters=50) for accuracy-first retrains, with LR for quick regression checks and pipeline canaries.

Spark profile sensitivity: Scaling the container-level Spark profile from 2c-2g \rightarrow 4c-4g \rightarrow 4c-6g reduces time (not AUC) predictably: RF and GBT benefit noticeably from more cores/GB, while LR gains are modest due to its low computational intensity. For a laptop-class setup, 4c-4g is a pragmatic default; 4c-6g provides headroom when data or HPO workload increases.

6.3 Reliability, governance, and safety

We codified guardrails so model evolution remains safe:

- **Registry gates:** A candidate is promoted only if it beats the incumbent AUC by $\delta \geq 0.005$ on the fixed evaluation split and shows no class-specific degradation in precision/recall. Manual approval is supported before Production.
- **Drift monitors:** PSI is computed on sentinel features; warning at > 0.2 , critical at > 0.3 triggers alerts and can schedule retraining. JSON reports are logged to MLflow and stored in MinIO for auditability.

- **Operational health.** /health and /reload endpoints support liveness checks and quick roll-forward to the newest Production model. The design supports canary releases and rollbacks by pinning image + model digests.
- **Reproducibility & audit.** Every run ties artifacts to parameters and code; environment snapshots round out the provenance chain. DVC ensures raw data lineage is explicit.

6.4 Limitations

Despite strong results, key limitations remain:

- **Single-split evaluation:** Fixed train/validation splits simplify comparison but may over- or under-estimate generalization. K-fold CV or repeated hold-outs would reduce variance.
- **Small, static dataset:** Titanic is purposefully compact; while we stress-tested the platform, scalability and failure modes should be re-validated on larger, messier, and potentially imbalanced datasets.
- **Container-level profiling:** Our resource charts capture end-to-end container usage. For deeper diagnosis (shuffle hotspots, skew), **Spark event logs** and History Server analytics should be fully integrated and summarized in the report.
- **Security hardening:** Secrets are injectable via env but should move to a central secrets manager for production environments; network encryption/TLS termination should be standardized for MinIO/MLflow/API.

6.5 Future work

1. **Stronger validation:** Introduce k-fold cross-validation and confidence intervals around AUC; add calibration (isotonic/Platt) and precision-recall curves to improve threshold selection and communication of risk.
2. **Spark-native profiling:** Enable event logging by default; parse stage/executor metrics (shuffle I/O, spill, skew) into tables and plots; correlate with container CPU/RSS to pinpoint bottlenecks.
3. **Observability & SLOs:** Add Prometheus metrics for the API (latencies, 5xx, request mix) and create Grafana dashboards; define p50/p95 latency and error-rate SLOs with alerting.

4. **Continuous training (CT):** Nightly retrains gated by offline metrics and drift; use Bayesian optimization/Hyperband to replace grid search for faster HPO at equal quality.
5. **Feature store & contracts:** Introduce a feature store (offline/online parity) and schema contracts with data quality tests (e.g., Great Expectations) to catch upstream breaks.
6. **Rollout safety:** Automate canary/blue-green flows with traffic shaping and rollbacks based on pre-declared guardrails; include shadow deployments for higher-risk changes.
7. **Scale & portability:** Migrate training orchestration to Kubernetes with a dedicated Spark operator; adopt object-storage lifecycle rules and retention policies for cost control at scale.

6.6 Summary

We set out to build a credible, end-to-end MLOps system that balances scientific rigor with operational practicality. The resulting platform is not just a model that performs well on Titanic; it is a repeatable process for moving data-driven ideas from notebooks to observable, governed, and maintainable services. The architecture stands on widely adopted components—Spark, MLflow, MinIO, FastAPI—integrated to preserve lineage, enforce gates, and keep day-2 operations simple.

Empirically, GBT is the accuracy leader on this task, while LR remains indispensable for speed and guardrails; RF provides a strong middle ground with excellent memory characteristics. With drift detection, automated retraining, and clear promotion criteria in place, the system is positioned to adapt as data or requirements evolve. Implementing the next tranche of work—CV, Spark-native profiling, observability, and safer rollouts—will convert this robust submission into a production-grade capability that remains sustainable as datasets grow, teams expand, and expectations

References:

- [1] **M. Zaharia, R. S. Xin, P. Wendell, T. Das, M. Armbrust, A. Dave, X. Meng, J. Rosen, S. Venkataraman, M. J. Franklin, A. Ghodsi, J. Gonzalez, S. Shenker, and I. Stoica**, “Apache Spark: A unified engine for big data processing,” *Communications of the ACM*, vol. 59, no. 11, pp. 56–65, Nov. 2016, doi: **10.1145/2934664**.
- [2] **A. Chen, A. Chow, A. Davidson, A. D’Cunha, A. Ghodsi, S. A. Hong, A. Konwinski, C. Mewald, S. Murching, T. Nykodym, P. Ogilvie, M. Parkhe, A. Singh, F. Xie, M. Zaharia, R. Zang, J. Zheng, and C. Zumar**, “Developments in MLflow: A system to accelerate the machine learning lifecycle,” in *Proc. Int. Workshop on Data Management for End-to-End Machine Learning (DEEM’20)*, 2020, doi: **10.1145/3399579.3399867**.
- [3] **N. del Río, B. Braun, M. Mirman, A. Guzman, M. Elhemali, M. Cafarella, M. Richardson, and B. Recht**, “Optimizing data pipelines for machine learning in feature stores,” *Proceedings of the VLDB Endowment*, vol. 16, no. 12, 2023, doi: **10.14778/3625054.3625060**.
- [4] **D. Petrović, J. Wong, and J. Rellermeier**, “Hopsworks feature store: A data platform for machine learning,” in *Companion of the 2024 Int’l Conf. on Management of Data (SIGMOD Companion)*, 2024, doi: **10.1145/3626246.3653389**.
- [5] **D. Franciasques, D. Smirnov, K. Ognawala, and S. Wagner**, “Machine learning experiment management tools: A mixed-methods study,” *Empirical Software Engineering*, 2024, doi: **10.1007/s10664-024-10444-w**.
- [6] **T. Hinder, E. Vaquet, and B. Hammer**, “One or two things we know about concept drift—Part A: theory, taxonomies and overall landscape,” *Frontiers in Artificial Intelligence*, vol. 7, 2024, doi: **10.3389/frai.2024.1330257**.
- [7] **T. Hinder, E. Vaquet, and B. Hammer**, “One or two things we know about concept drift—Part B: application-driven case studies,” *Frontiers in Artificial Intelligence*, vol. 7, 2024, doi: **10.3389/frai.2024.1330258**.
- [8] **M. R. Haas and L. Sibbald**, “Measuring data drift with the Unstable Population Indicator,” *Data Science*, 2024, doi: **10.3233/DS-240059**.
- [9] **A. Morales-Hernández, A. Fernández, and F. Herrera**, “A survey on multi-objective hyperparameter optimization: Current trends and open challenges,” *Artificial Intelligence Review*, vol. 56, 2023, doi: **10.1007/s10462-022-10359-2**.
- [10] **Iterative, Inc.**, “Data Version Control (DVC) v3.x,” *Zenodo*, 2024, doi: **10.5281/zenodo.11266562**.

